

Mod 5 - Lecture 9:

Multi-Level Indexes



- Deepak Poonia

(IISc Bangalore; GATE AIR 53 & 67)

Content of this Lecture:

1. Multilevel Indices
 2. ISAM
 3. Practice + GATE Questions
- DBMS Complete Course Link:

<https://www.goclasses.in/courses/Database-Management-Systems>

Instructor:
Deepak Poonia
IISc Bangalore

GATE CSE AIR 53; AIR 67; AIR 206; AIR 256

DBMS Complete Course Link:

<https://www.goclasses.in/courses/Database-Management-Systems>



GATE 2024 COMPLETE COURSE CS-IT

(1 YEAR + 3 MONTHS)



GATE 2024 COMPLETE
COURSE CS - IT
(1 YEAR + 3 MONTHS)



GATE 2023 COMPLETE COURSE CS-IT

(6 MONTHS)



GATE 2023 COMPLETE
COURSE CS-IT
(6 MONTHS)



GATE 2025 COMPLETE COURSE CS-IT

(2 YEARS)



GATE 2025 COMPLETE
COURSE CS - IT
(2 YEARS + 3 MONTHS)



Discrete Mathematics



2023 Discrete Mathematics

★★★★★ 5.0 (62 ratings)

Deepak Poonia (MTech IISc Bangalore)

Free



C Programming



2023 C Programming

★★★★★ 5.0 (59 ratings)

Sachin Mittal (MTech IISc Bangalore)

Free



www.goclasses.in

Discrete Mathematics

2023 Discrete Mathematics

Learn, Understand, Discuss. "GO" for the Best.

★★★★★ 5.0 (62 ratings)

5997 Learners Enrolled

Language: English

Instructors: Deepak Poonia (MTech IISc Bangalore, GATE CSE AIR 53; 67)

FREE

On
“GATE-
Overflow

”
Website

G GO CLASSES + Data

G GO CLASSES



**GO Test Series
is now**

**GATE Overflow + GO Classes
2-IN-1 TEST SERIES**

**Most Awaited
GO Test Series
is Here**

REGISTER NOW

<http://tests.gatecse.in/>

100+ More than 100 Quality Tests.

15 Mock Tests.

FROM

14th April

+91 - 6302536274 +91 9499453136





GATE 2023

LIVE

Live + Recorded Lectures

Daily Home Work + Solution

Watch Any Time + Any Number of Times

Summary Lectures For Every Topic

Practice Sets From Standard Resources

**Enroll Now**

+91 - 6302536274

www.goclasses.in



linkedin.com/company/go-classes



instagram.com/goclasses_cs



Join **GO+ GO Classes Combined Test Series** for **BEST** quality tests, matching GATE CSE Level:

Visit www.gateoverflow.in website to join Test Series.

1. **Quality Questions:** No Ambiguity in Questions, All Well-framed questions.
2. Correct, **Detailed Explanation**, Covering Variations of questions.
3. **Video Solutions.**

GO Test Series Available Now
Revision Course
GATE PYQs Video Solutions

SPECIAL



EXPLORE OUR FREE COURSES

Free Discrete Mathematics Complete Course
C-Programming Complete Course



Best Mentorship and Support



Sachin Mittal
(CO-Founder GOCLASSES)

MTech IISc Bangalore
Ex Amazon scientist
GATE AIR 33

Deepak Poonia
(CO-Founder GOCLASSES)

MTech IISc Bangalore
GATE AIR 53; 67

Dr. Arjun Suresh
(Mentor)

Founder GATE Overflow
Ph.D. INRIA France
ME IISc Bangalore
Post-doc The Ohio State University

www.goclasses.in

GATE CSE 2023

(LIVE + RECORDED COURSE)

NO PREREQUISITES
FROM BASICS, IN - DEPTH

Enroll Now

Quality Learning.

Weekly Quizzes.

Summary Lectures.

Daily Homeworks
& Solutions.Interactive Classes
& Doubt Resolution.ALL GATE PYQs
Video Solutions.Doubts Resolution
by Faculties on Telegram.Selection Oriented
Preparation.Standard Resources
Practice Course.

+91- 6302536274

Visit Website for More Details

www.goclasses.in



NOTE :

Complete Discrete Mathematics & Complete C-Programming

Courses, by GO Classes, are **FREE** for ALL learners.

Visit here to watch : <https://www.goclasses.in/s/store/>

SignUp/Login on Goclasses website for free and start learning.



Download the GO Classes Android App:

<https://play.google.com/store/apps/details?id=com.goclasses.courses>

Search “GO Classes”
on Play Store.



www.goclasses.in

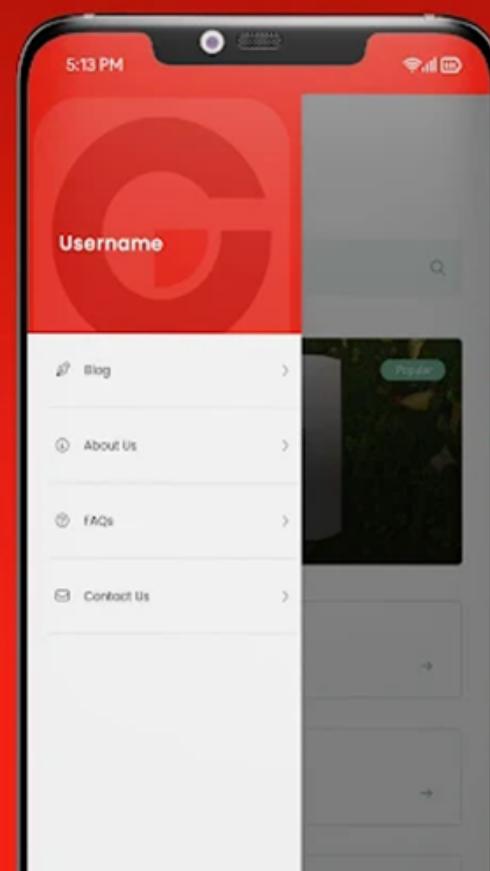
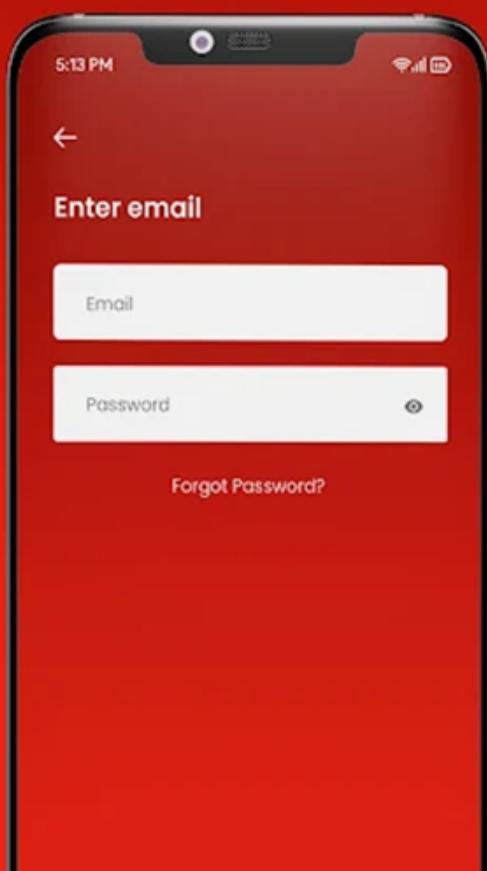
Hassle-free learning
On the go!

Gain expert knowledge



Continue with email

Don't have an account? [Sign up](#)





We are on **Telegram**. Contact Us for any help.

Join GO Classes **Resources**, Notes, Content, information **Telegram Channel**:

Public Username: **GOCLASSES_CSE**

Join GO Classes **Doubt Discussion** Telegram Group :

Username: **GATECSE_Goclasses**

(Any doubt related to Goclasses Courses can also be asked here.)

Join GATEOverflow **Doubt Discussion** Telegram Group :

Username: **gateoverflow_cse**

*Types of Single level
Ordered Indexes:*

Summary

Index Structures

Index: A disk data structure

- enables efficient retrieval of a record given the value (s) of certain attributes
 - indexing attributes

Primary Index:

Index built on *ordering key* field of a file

Clustering Index:

Index built on *ordering non-key* field  of a file

Secondary Index:

Index built on any *non-ordering* field of a file

17.1.4 Summary

To conclude this section, we summarize the discussion of index types in two tables. Table 17.1 shows the index field characteristics of each type of ordered single-level index discussed—primary, clustering, and secondary. Table 17.2 summarizes the properties of each type of index by comparing the number of index entries and specifying which indexes are dense and which use block anchors of the data file.



Table 17.1 Types of Indexes Based on the Properties of the Indexing Field

	Index Field Used for Physical Ordering of the File	Index Field Not Used for Physical Ordering of the File
Indexing field is key	Primary index	Secondary index (Key)
Indexing field is nonkey	Clustering index	Secondary index (NonKey)

Table 17.2 Properties of Index Types

Type of Index	Number of (First-Level) Index Entries	Dense or Nondense (Sparse)	Block Anchoring on the Data File
Primary	Number of blocks in data file	Nondense	Yes
Clustering	Number of distinct index field values	Nondense	Yes/no ^a
Secondary (key)	Number of records in data file	Dense	No
Secondary (nonkey)	Number of records ^b or number of distinct index field values ^c	Dense or Nondense	No

^aYes if every distinct value of the ordering field starts a new block; no otherwise.

^bFor option 1.

^cFor options 2 and 3.

Next Topic:

Ordered Indexes

Multilevel

So far:

Index is ordered

Single level ordered Indexes

I index blocks, B Data Blocks

PI

CI

SI

[on key]

[on Nonkey]

Access Cost (I/o cost): $\lceil \log_2 I \rceil + 1$

Binary Search
on Index

for
Data Block

I/o cost: $\lceil \log_2 I \rceil + 1$ binary search

To Reduce I/o Cost:

① Put Index in main memory

Problem 1: Index larger than main memory

Problem 2: Even if Index Size < mm size

then $\frac{I}{mm} \rightarrow$ mm has many things to accomodate
So mm may not be able to take index

I/o cost:

$$\lceil \log_2 I \rceil + 1$$

binary
search

Date
Block

To Reduce I/o cost:

① Put Index in main memory: $\Rightarrow \frac{\text{I/o cost}}{\text{Index}}$

Good only when Index size is very small



at most 2-3
Index Blocks

I/o cost: $\lceil \log_2 I \rceil + 1$

binary search

To Reduce I/o Cost:

① Put Index in main memory

↳ may be okay for Primary Index

small size

② Multilevel Index ✓

Reduce it



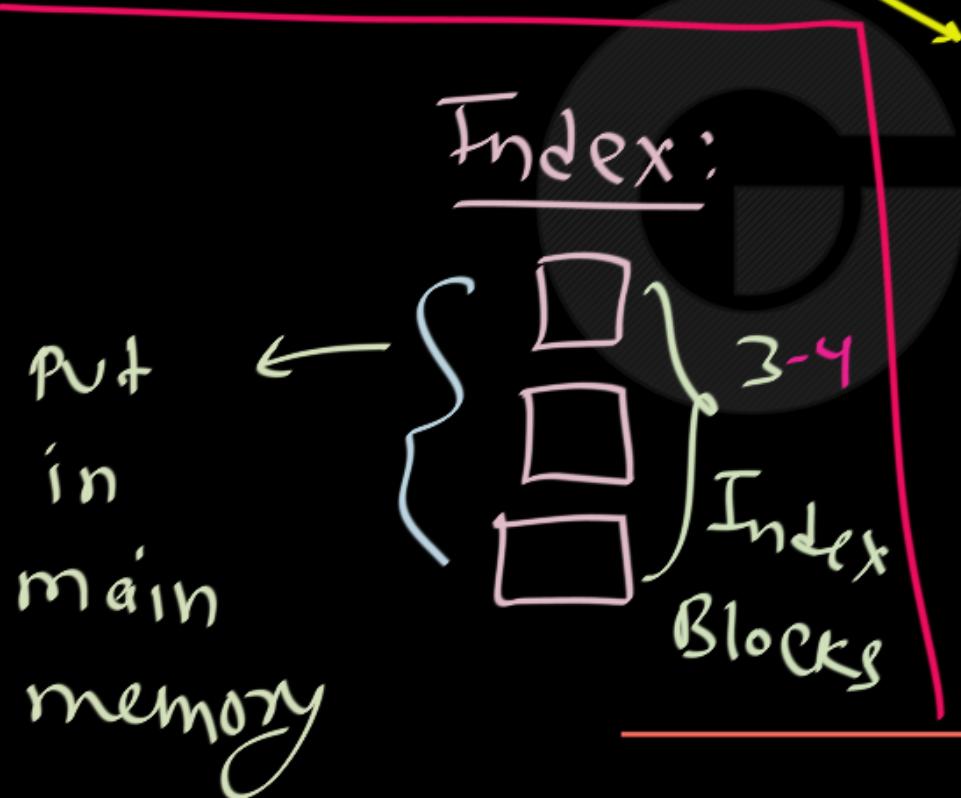
+ 1 → for data block

Q: we go for multilevel index because
'Index doesn't fit in one block ?'



Reason for multilevel index??

Q: we go for multilevel index because
'Index doesn't fit in one block ?'



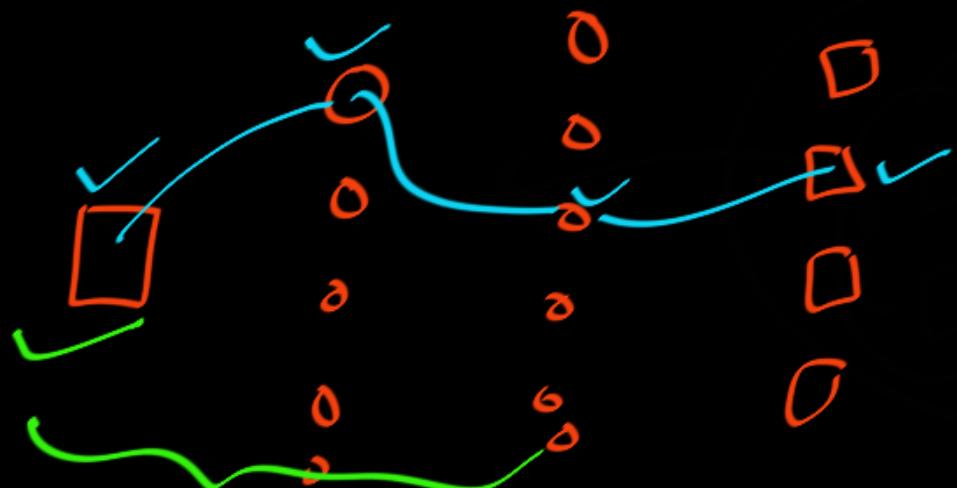
Reason for multilevel index??

No

No

Once we Decide to go for
multilevel Index then by default

Outer most level has one Block.



I/o cost = $4 = 3 + 1$

Index levels
Data Block

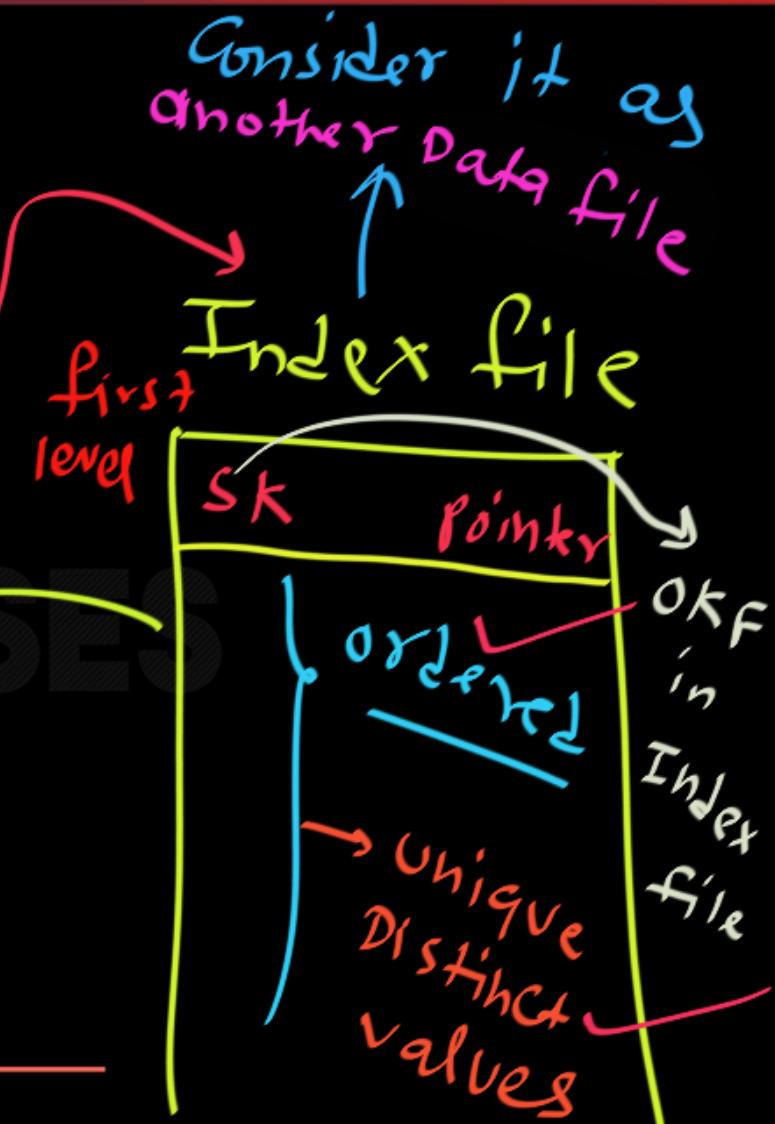
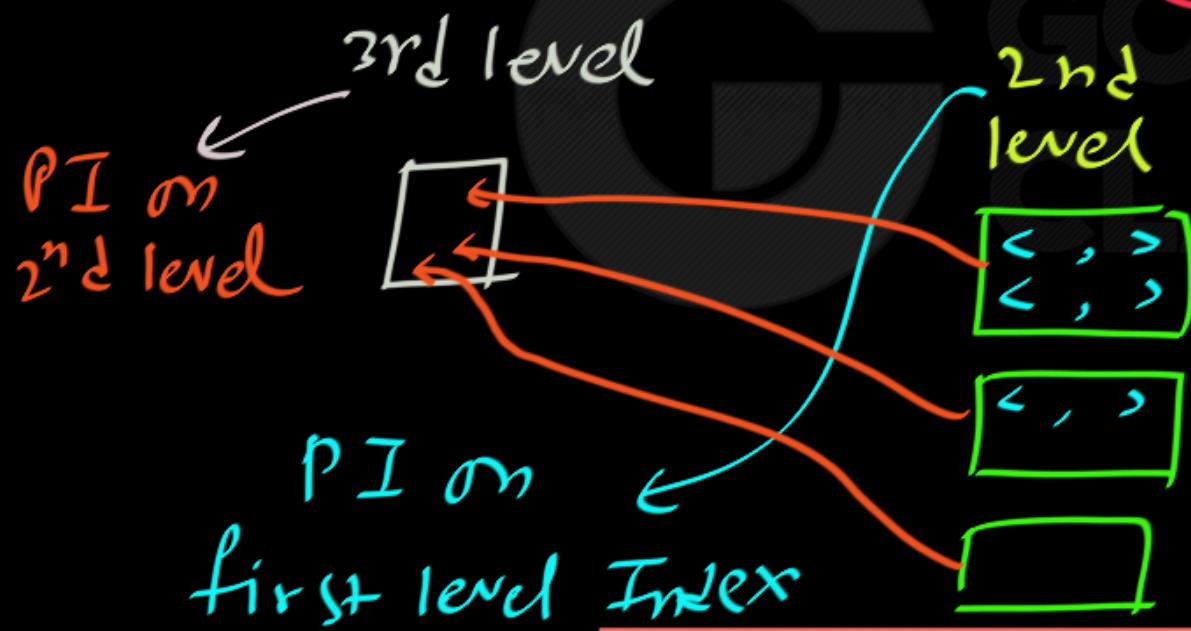
17.2 Multilevel Indexes

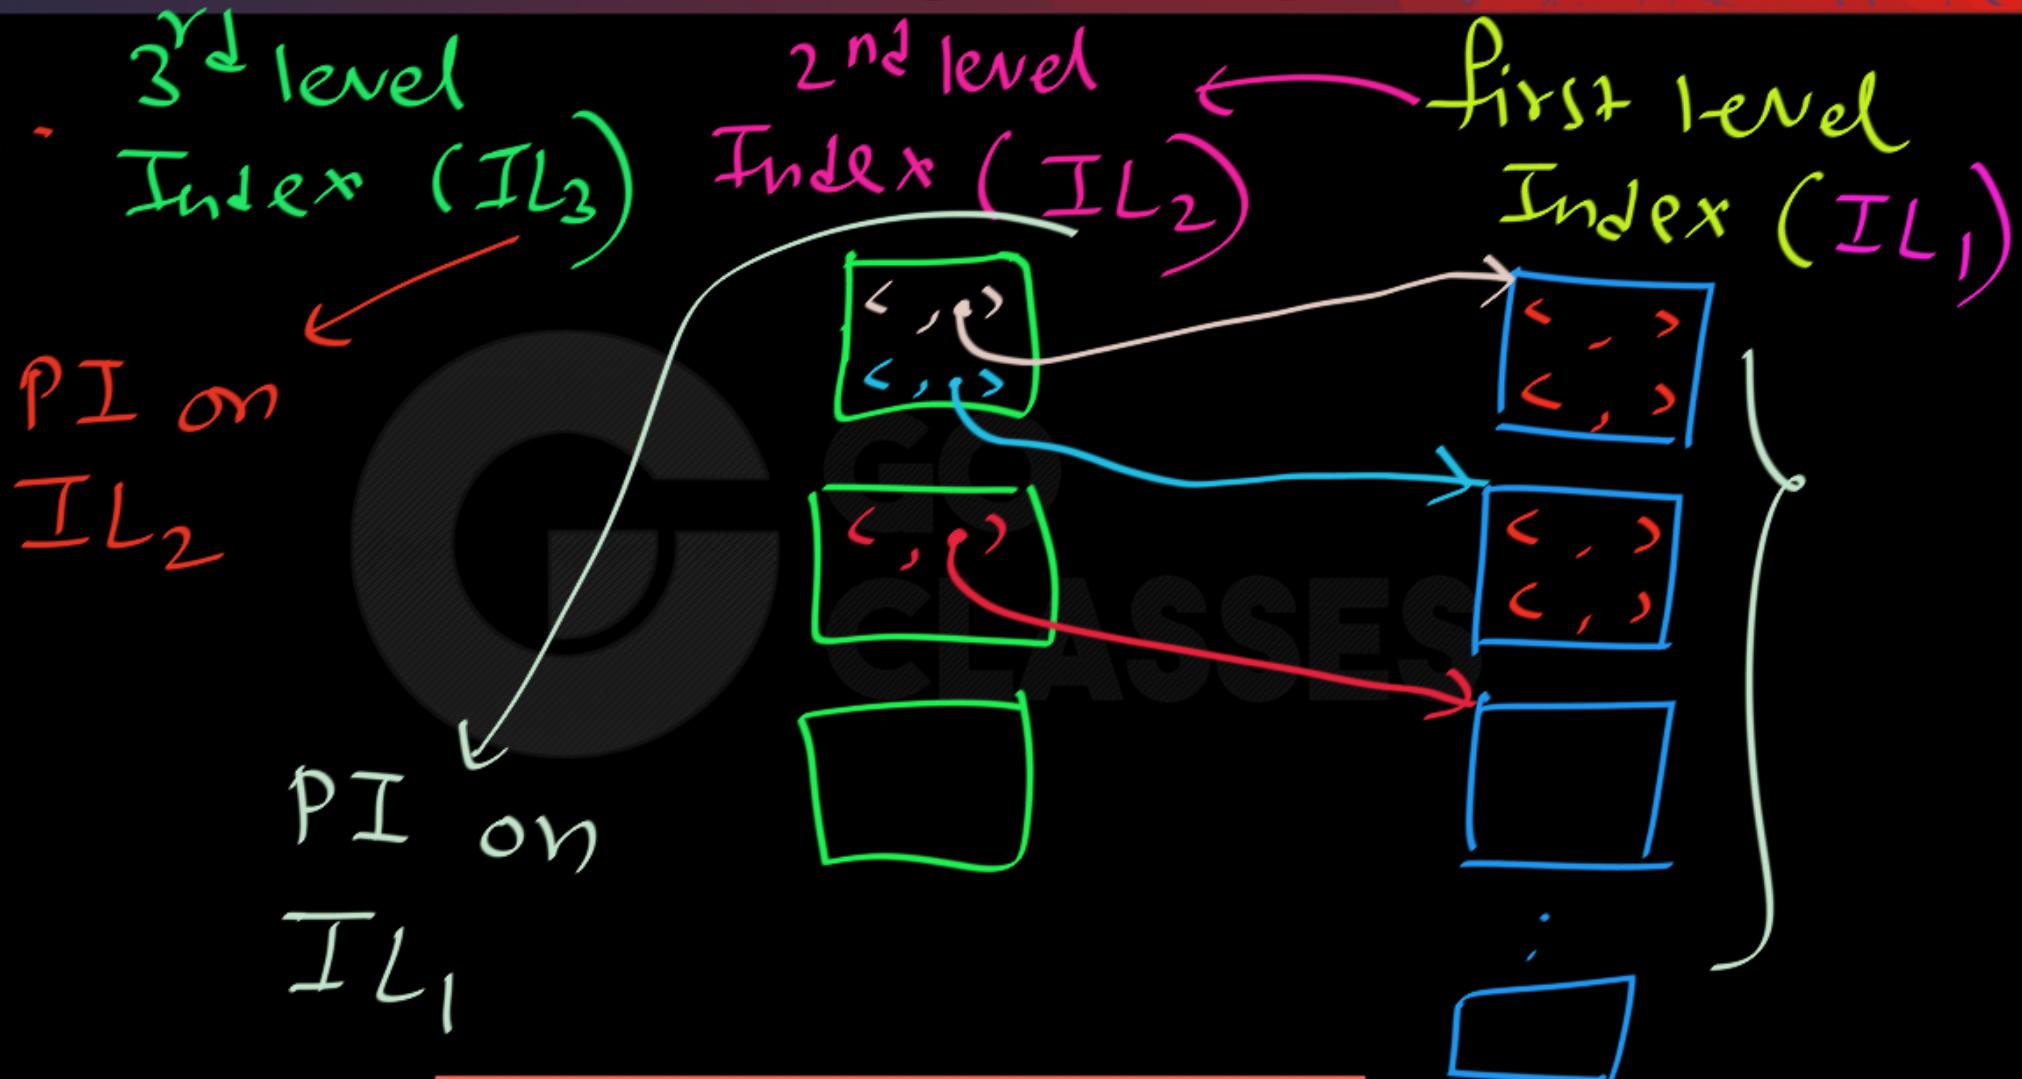
The indexing schemes we have described thus far involve an ordered index file. A binary search is applied to the index to locate pointers to a disk block or to a record (or records) in the file having a specific index field value. A binary search requires approximately $(\log_2 b_i)$ block accesses for an index with b_i blocks because each step of the algorithm reduces the part of the index file that we continue to search by a factor of 2. This is why we take the log function to the base 2.

If an index is small enough to be kept entirely in main memory, the search time to find an entry is low. However, if the index is so large that not all of it can be kept in memory, index blocks must be fetched from disk when required. (Even if an index is smaller than the main memory of a computer, main memory is also required for a number of other tasks, so it may not be possible to keep the entire index in memory.) The search for an entry in the index then requires several disk-block reads.

Single level Indexes:

Primary Index, CI, SI





MLI : (multilevel Index) :

first
level Index

IL₁

normal
PI / CI / SI
on

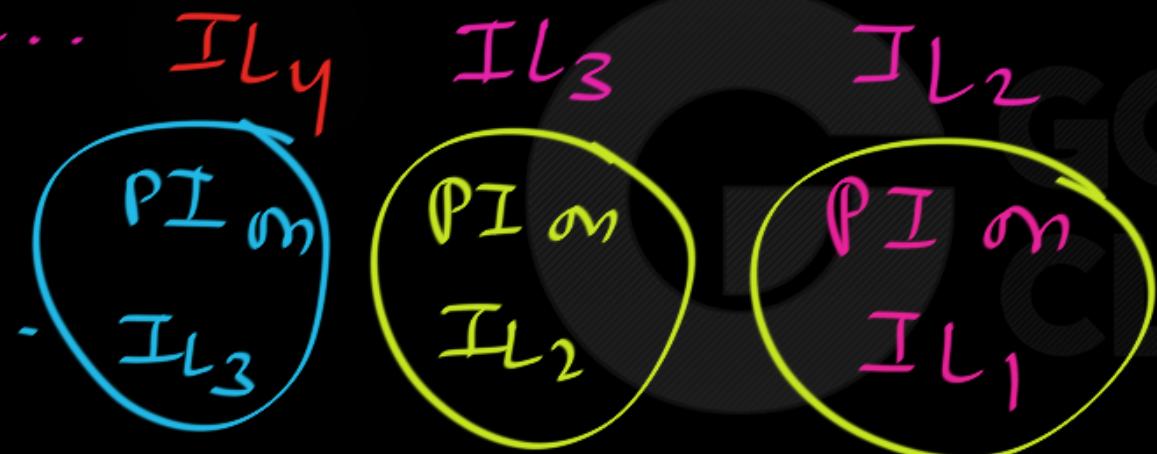
Datefile

Date file



b

b/
locky



MLI : (multilevel Index) ;

Note : 2nd level and greater are

ALWAYS

SPARSE

ALWAYS

MLI : (multilevel Index) ;

Note : 2nd level and greater are
Always SPARSE

Note : Index level K is PI (sparse)
on Index level K-1.

MLI : (multilevel Index) ;

Note : 2nd level and greater are

Always SPARSE

Note : Index level 1 is PI / CI / SI Normal
on Data file

MLI : (multilevel Index) ;

Note : 2nd level and greater are
Always SPARSE

Note : Last (outermost) level is
called Top level [Inner most level
is 1st level]

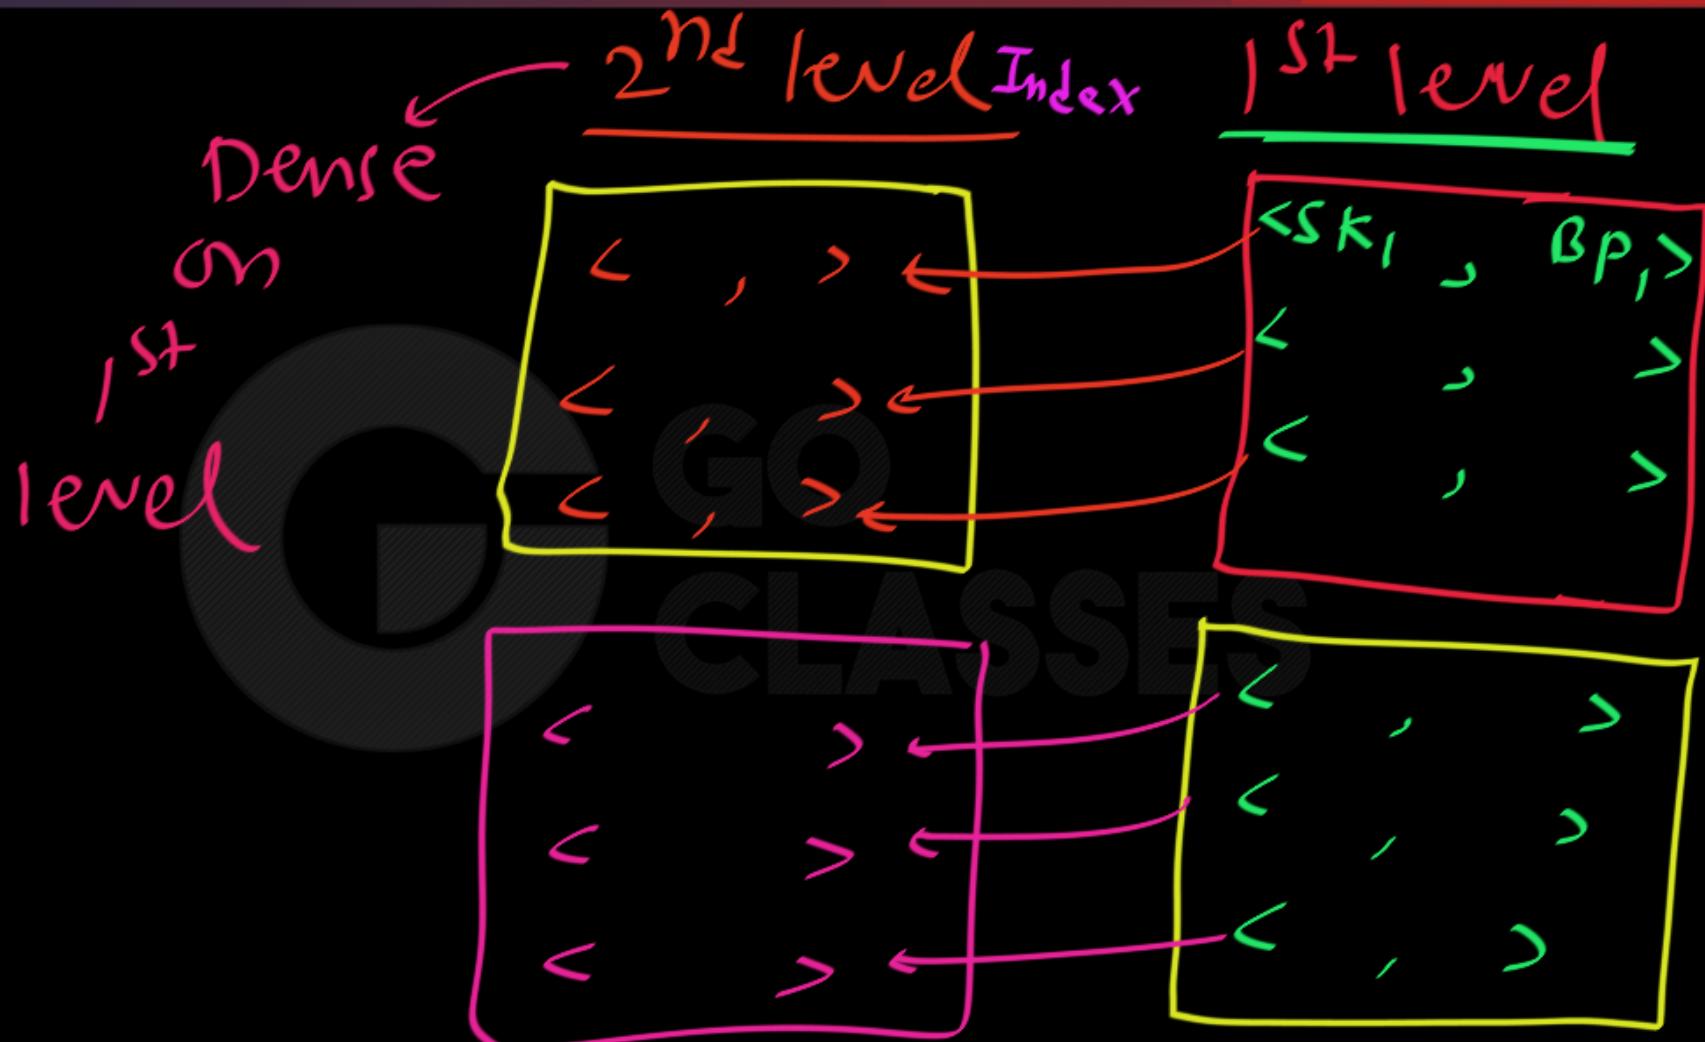
Q: Assume 1st level is PI on Datafile
8 2nd level is Dense index on 1st level then

Problem ??.

Q: Assume 1st level is P I on Datafile
8 2nd level is Dense index on 1st
level then

Problem ??

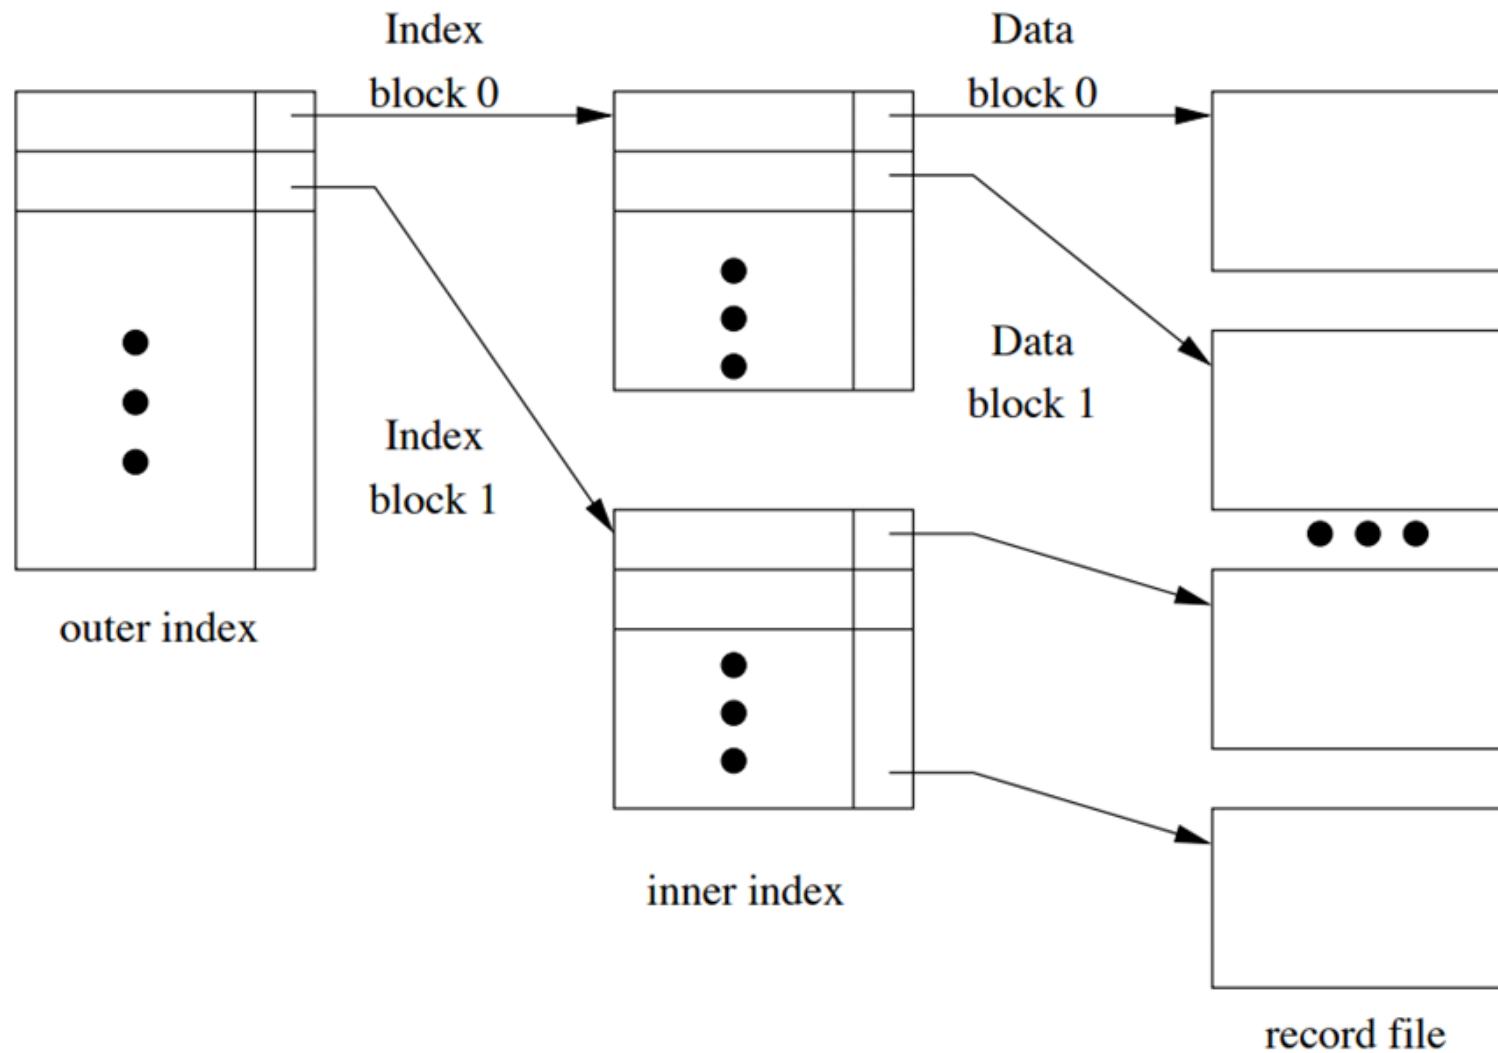
All Index levels
Copy on each other



Multi-Level Index

- If primary index does not fit in memory, access to records becomes expensive
- To reduce number of disk accesses to index entries, treat primary index on disk as sequential file and construct a sparse index on it.
 - outer index → a sparse index of primary index
 - inner index → the primary index file
- If even outer index is too large to fit in main memory, yet another level of index can be created, and so on.

- Multilevel Index structure



Making the Index Multi-level

Index file – itself an ordered file
– another level of index can be built

Multilevel Index –

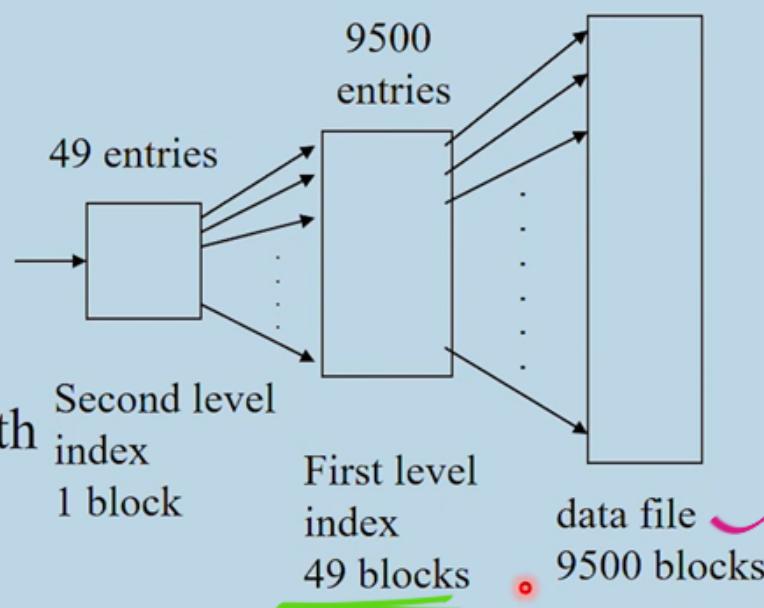
Successive levels of indices are built till the last level has one block

height – no. of levels

block accesses: height + 1
(no binary search required)

For the example data file:

No of block accesses required with
multi-level primary index: 3
without any index: 14



φ : If Data file is ordered : # Data Blocks

① Approach 1 : without index :

$$I/O \text{ cost} = \lceil \log_2 9500 \rceil = 14$$

② App 2 : with single level index :

$$I/O \text{ cost} = \lceil \log_2 I \rceil + 1 = 6 + 1 = 7$$

③ App3 : mLI : $2 + 1 = 3$

Note: MLI is better choice than
Binary Search on Single level.



FUNDAMENTALS OF Database Systems

7th Edition Ramez Elmasri, Shamkant B. Navathe*Exercise 17.18**(Multi part question)*

Q 17.18: FUNDAMENTALS OF Database Systems 7th Ed, Navathe, Elmasri

17.18. Consider a disk with block size $B = 512$ bytes. A block pointer is $P = 6$ bytes long, and a record pointer is $P_R = 7$ bytes long. A file has $r = 30,000$ EMPLOYEE records of fixed length. Each record has the following fields: Name (30 bytes), Ssn (9 bytes), Department_code (9 bytes), Address (40 bytes), Phone (10 bytes), Birth_date (8 bytes), Sex (1 byte), Job_code (4 bytes), and Salary (4 bytes, real number). An additional byte is used as a deletion marker.

- a. Calculate the record size R in bytes. → record header
- b. Calculate the blocking factor bfr and the number of file blocks b , assuming an unspanned organization.

Block Size: $b = \underline{512 \text{ Bytes}} = 512 \text{ B}$

Block pointer = $p = 6 \text{ Bytes} = 6 \text{ B}$

Record $r_r = p_R = 7 \text{ B}$

#file Records = $\gamma = 30,000$

Record Size = $30 + 9 + 9 + 40 + 10 + 8 + 1 + 4 + 4 + 1$
= 116 B (Fixed length)

Blocking factor of file : bfr :

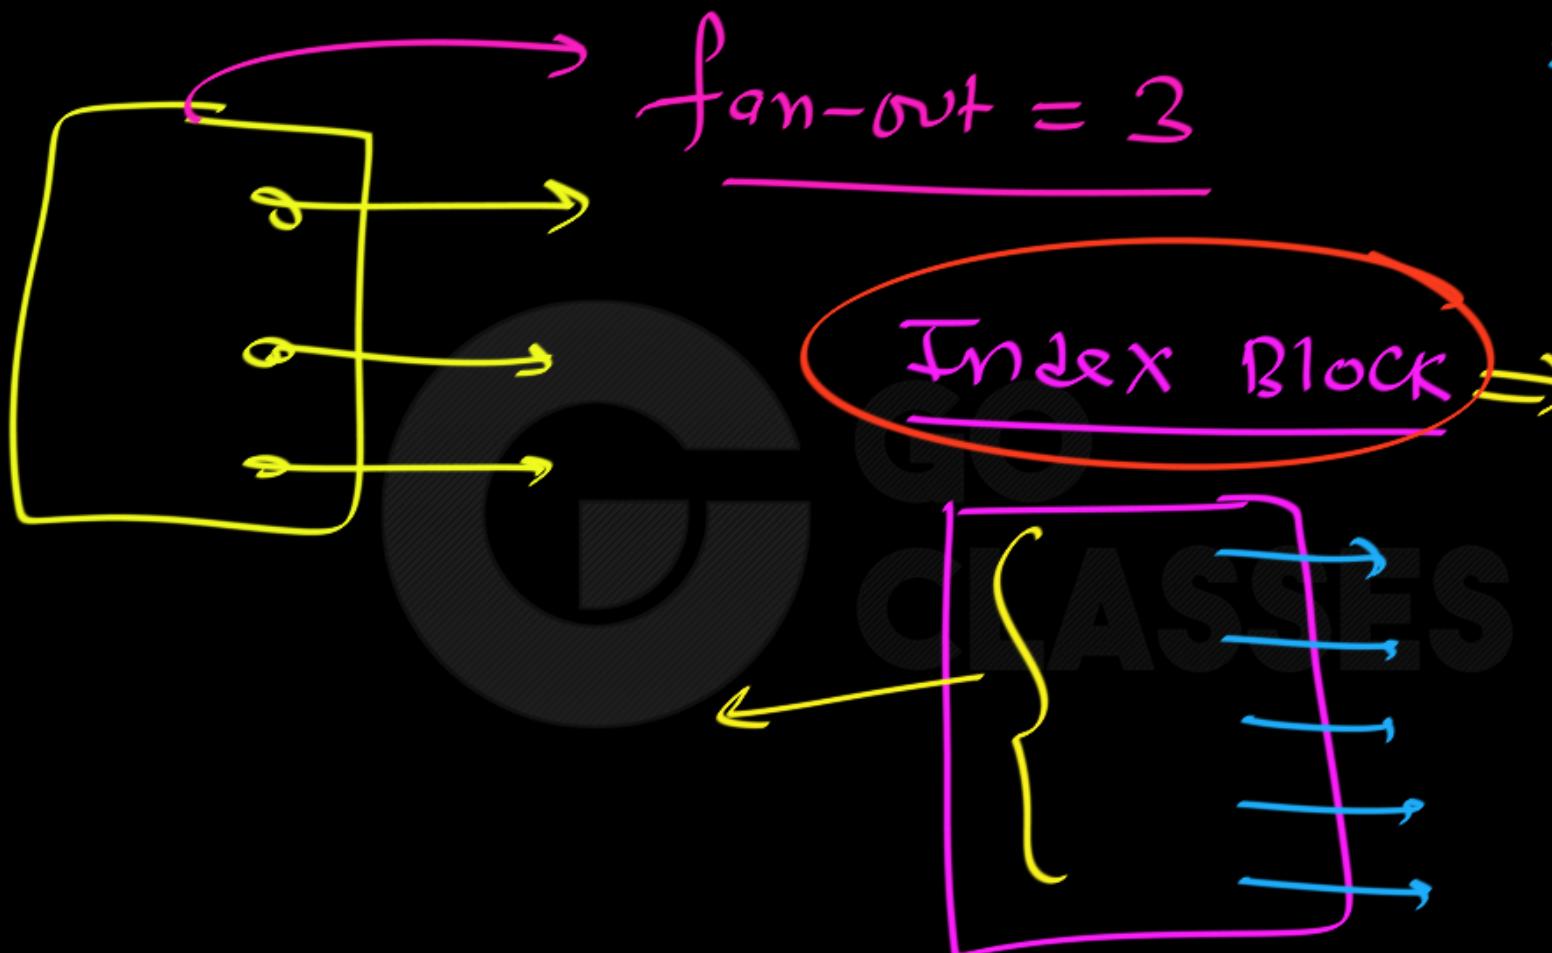
#Records / Block

#File Blocks needed :

$$\left\lceil \frac{512 \text{ B}}{114 \text{ B}} \right\rceil = 4$$

$$\left\lceil \frac{30,000}{4} \right\rceil = 7500 \text{ Blocks}$$

in
one
block



Index Records
per Block

Blocking Factor

\equiv fan-out

$\equiv f_o$

Q 17.18: FUNDAMENTALS OF Database Systems 7th Ed, Navathe, Elmasri

17.18. Consider a disk with block size $B = 512$ bytes. A block pointer is $P = 6$ bytes long, and a record pointer is $P_R = 7$ bytes long. A file has $r = 30,000$ EMPLOYEE records of *fixed length*. Each record has the following fields: Name (30 bytes), Ssn (9 bytes), Department_code (9 bytes), Address (40 bytes), Phone (10 bytes), Birth_date (8 bytes), Sex (1 byte), Job_code (4 bytes), and Salary (4 bytes, real number). An additional byte is used as a deletion marker.

- c. Suppose that the file is ordered by the key field Ssn and we want to construct a *primary index* on Ssn. Calculate (i) the index blocking factor bfr_i (which is also the index fan-out fo); (ii) the number of first-level index entries and the number of first-level index blocks; (iii) the number of levels needed if we make it into a multilevel index; (iv) the total number of blocks required by the multilevel index; and (v) the number of block accesses needed to search for and retrieve a record from the file—given its Ssn value—using the primary index.

Block size: $b = 512 \text{ Bytes} = \underline{512 \text{ B}}$

Block pointer = $p = 6 \text{ Bytes} = \underline{6 \text{ B}}$

Record $r_r = p_R = \underline{7 \text{ B}}$

#file Records = $\gamma = 30,000$

Record size = $30 + 9 + 9 + 40 + 10 + 8 + 1 + 4 + 4 + 1$
= 116 B (fixed length)

File Blocking factor = 4; #file Blocks = 7500

Second level

IndexRecord size

$$= 9 + 6 = 15 \text{ B}$$

sk size

BP
size

SSN

Blocking factor of IL2 = 34

3rd level

$$\# IL_3 \text{ entries} = 7 \checkmark$$

Blocking factor of IL_3 = 34

2nd level

$$\# IL_2 \text{ Entries} = 221$$

$\# IL_2 \text{ Blocks}$

$$= \left\lceil \frac{221}{34} \right\rceil = 7$$

1st level

$$\# IL_1 \text{ entries} = 7500 \checkmark$$

$$\# IL_1 \text{ blocks} = \left\lceil \frac{7500}{34} \right\rceil = 221$$

- (i) 34
 (ii) 7500 ; 221
 (iii) 3 Index level
 (iv) $1 + 7 + 221 = 229$

⑤ Search on SSN: I/o cost = $3 + 1 = 4$ ✓
Data Block
Search on Address = I/o cost \Rightarrow Avg = $\frac{7500}{2} = 3750$ ✓

Q 17.18: FUNDAMENTALS OF Database Systems 7th Ed, Navathe, Elmasri

- 17.18. Consider a disk with block size $B = 512$ bytes. A block pointer is $P = 6$ bytes long, and a record pointer is $P_R = 7$ bytes long. A file has $r = 30,000$ EMPLOYEE records of *fixed length*. Each record has the following fields: Name (30 bytes), Ssn (9 bytes), Department_code (9 bytes), Address (40 bytes), Phone (10 bytes), Birth_date (8 bytes), Sex (1 byte), Job_code (4 bytes), and Salary (4 bytes, real number). An additional byte is used as a deletion marker.
- d. Suppose that the file is not ordered by the key field Ssn and we want to construct a secondary index on Ssn. Repeat the previous exercise (part c) for the secondary index and compare with the primary index.

Block size: $b = 512 \text{ Bytes} = 512 \text{ B}$

Block pointer = $p = 6 \text{ Bytes} = 6 \text{ B}$

Record $r_i = p_R = 7 \text{ B}$

#file Records = $\gamma = 30,000$

Record size = $30 + 9 + 9 + 40 + 10 + 8 + 1 + 4 + 4 + 1$
= 116 B (fixed length)

File Blocking factor = 4; #file Blocks = 7500

① Index Entry Size :

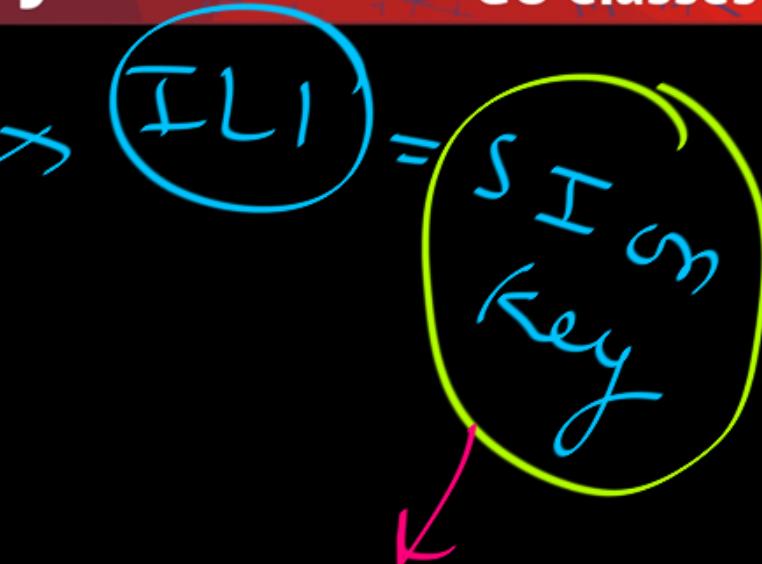
first level index : \Rightarrow

$$9 + 7 = 16 \text{ B}$$

ssn record pointer

2nd & Higher levels

$$\text{Index Record Size} = 9 + 6 = 15$$



Record pointer
OR
Block pointer

(2)

Blocking factor of IL1 :

$$\left\lfloor \frac{512B}{16B} \right\rfloor = 32$$

Blocking factor

Index records
per block

$$\left\lfloor \frac{512B}{15B} \right\rfloor = \underline{\underline{34}}$$

of IL2 & Higher levels:

③ #first level Entries : #Records in file
Index = 30,000

SI m key

ILI Blocks =

$$\left\lceil \frac{30,000}{32} \right\rceil = 938$$

③ # 2ⁿ⁻¹ level Entries : # blocks on IL₁
↓ Index = 938

$$\frac{\# \text{IL}_2 \text{ Blocks}}{3^4} = \left\lceil \frac{938}{3^4} \right\rceil = 28$$

3ⁿ level index entries = # blocks on 2ⁿ⁻¹ level
= 28

Single Block

level in MLI = 3

Index Blocks (total) = $1 + 28 + 938$
= 967

- ④ Access Cost:
Search on ssn: $3 + 1 \rightarrow$ Data Block
= 4 ✓

14.1.3 Secondary Indexes

- A secondary index is also an ordered file with two fields.
- The first field of the index is of the same data type as some *nonordering field* of the data file that is an indexing field.
- The second field of the index is either a block pointer or a record pointer.

(170 points) Consider a disk with block size $B=512$ bytes. A block pointer is $P = 6$ bytes long, and a record pointer is $PR = 7$ bytes long. A file has $r = 30,000$ EMPLOYEE records of fixed length. Each record has the following fields: NAME (30 bytes), SSN (9 bytes), DEPARTMENTCODE (9 bytes), ADDRESS (40 bytes), PHONE (9 bytes), BIRTHDATE (8 bytes), SEX (1 byte), JOBCODE (4 bytes), SALARY (5 bytes, real number).

4. (50 points) Suppose the file is **not ordered** by the key field SSN and we want to construct a secondary index on SSN. Repeat the previous exercise (part 3) for the secondary index and compare with the primary index. Assumption is the first level indexes use the record pointers and the higher-level indexes use the block pointers.
- a) the index blocking factor bfr (which is also the index fan-out fo)
 - b) the number of first-level index entries and the number of first-level index blocks
 - c) the number of levels needed if we make it into a multi-level index.
 - d) the total number of blocks required by the multi-level index.
 - e) the number of block accesses needed to search for and retrieve a record from the file--given its SSN value--using the **secondary** index.

Q 17.18: FUNDAMENTALS OF Database Systems 7th Ed, Navathe, Elmasri

- 17.18. Consider a disk with block size $B = 512$ bytes. A block pointer is $P = 6$ bytes long, and a record pointer is $P_R = 7$ bytes long. A file has $r = 30,000$ EMPLOYEE records of *fixed length*. Each record has the following fields: Name (30 bytes), Ssn (9 bytes), Department_code (9 bytes), Address (40 bytes), Phone (10 bytes), Birth_date (8 bytes), Sex (1 byte), Job_code (4 bytes), and Salary (4 bytes, real number). An additional byte is used as a deletion marker.

CLASSES
Part e: On Next Page

e. Suppose that the file is *not ordered* by the nonkey field Department_code and we want to construct a *secondary index* on Department_code, using option 3 of Section 17.1.3, with an extra level of indirection that stores record pointers. Assume there are 1,000 distinct values of Department_code and that the EMPLOYEE records are evenly distributed among these values. Calculate (i) the index blocking factor bfr_i (which is also the index fan-out fo); (ii) the number of blocks needed by the level of indirection that stores record pointers; (iii) the number of first-level index entries and the number of first-level index blocks; (iv) the number of levels needed if we make it into a multilevel index; (v) the total number of blocks required by the multilevel index and the blocks used in the extra level of indirection; and (vi) the approximate number of block accesses needed to search for and retrieve all records in the file that have a specific Department_code value, using the index.

SI on Non-key

Three implementations

variable length
index records

I₁:

(sk, BP₁, BP₂, BP₃)

Not preferred

I₂:

20 BP₁
20 BP₂
20 BP₃

I₃:

by implementation

New level
of indexation

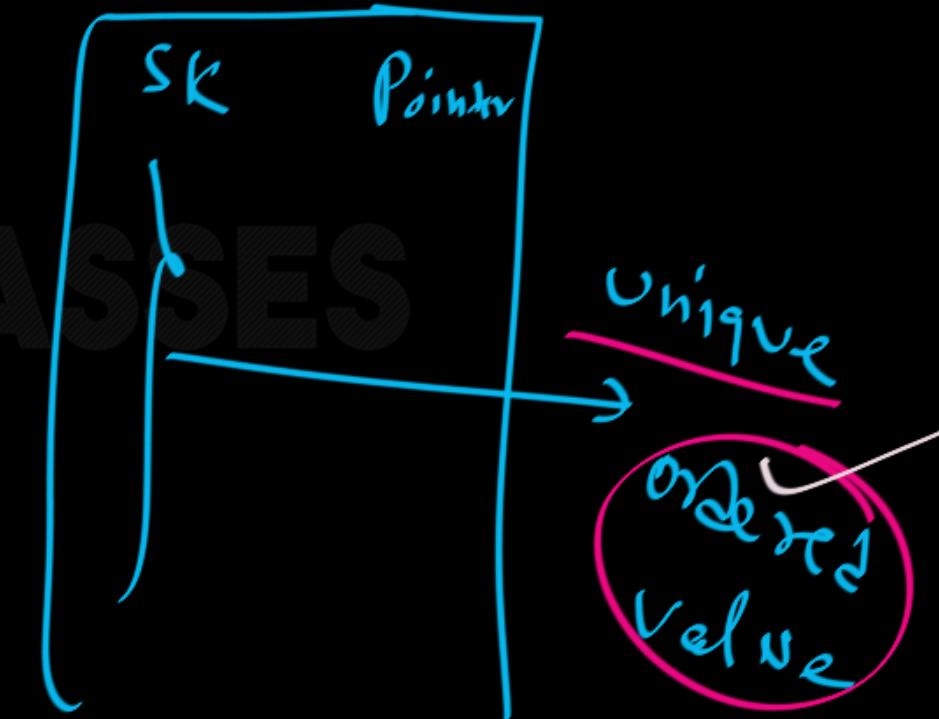


PI, CI, SI

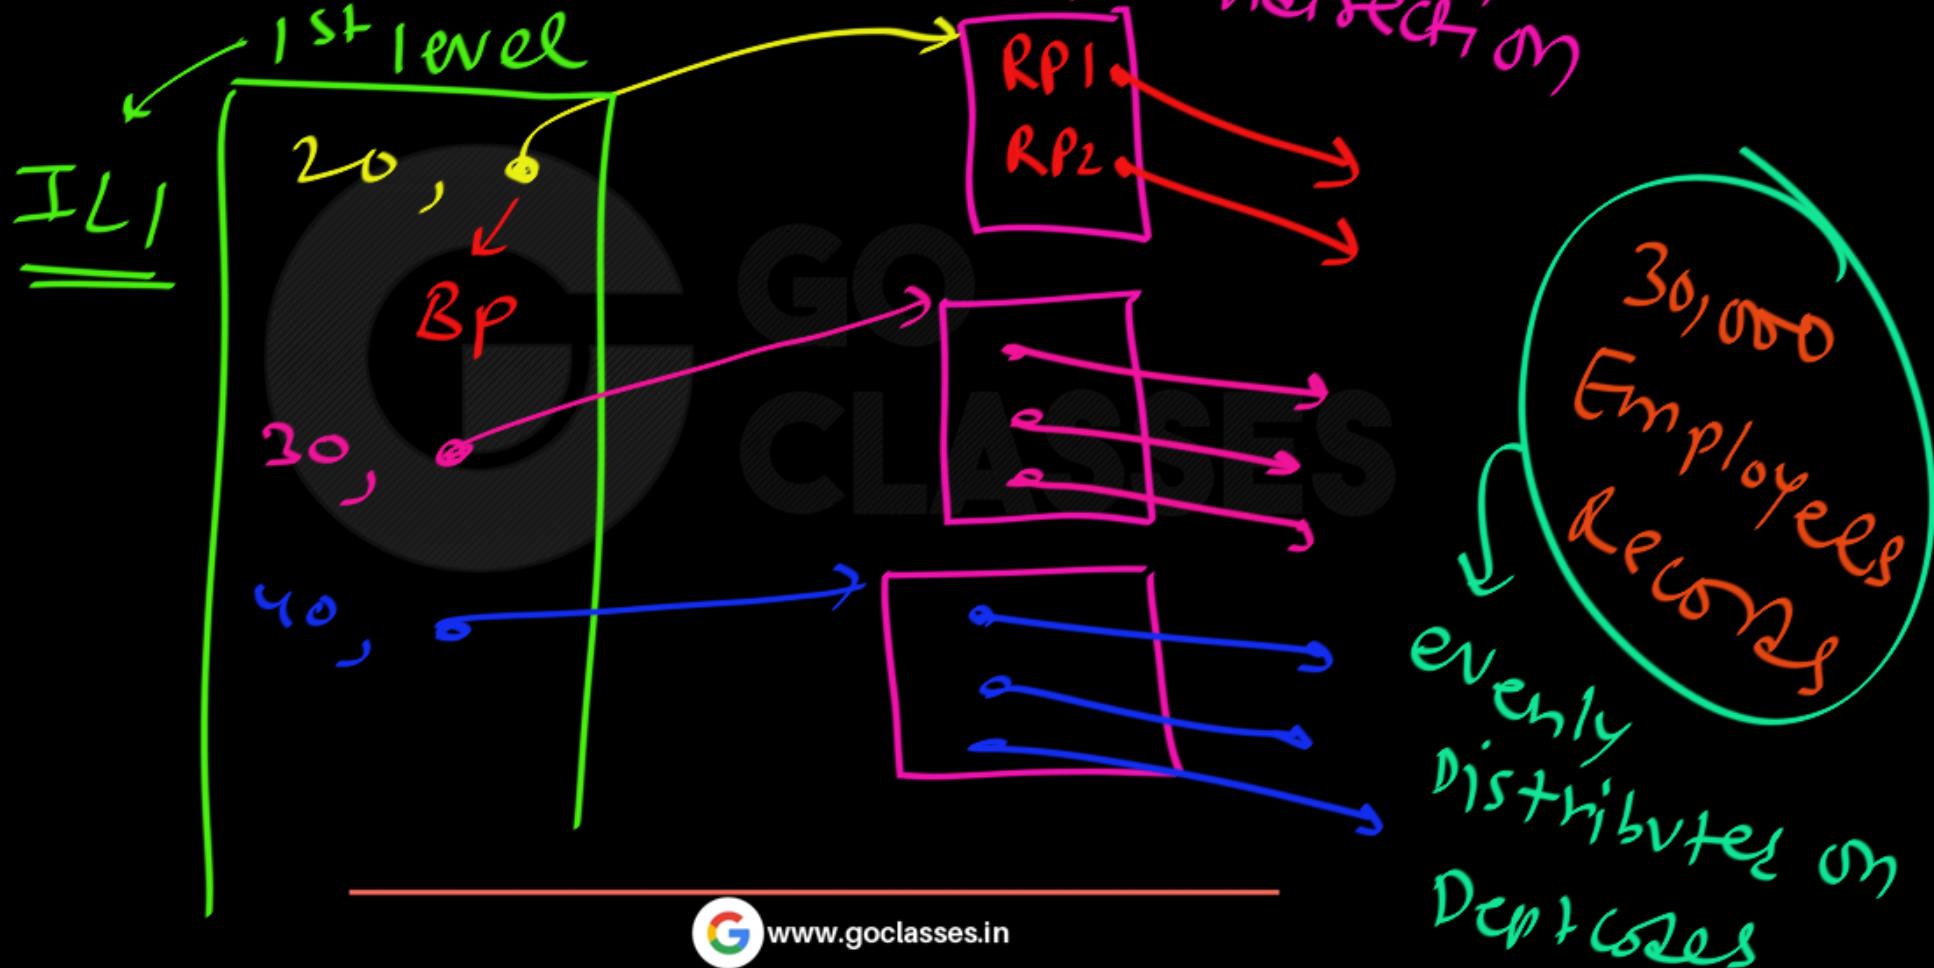
on 1st level :



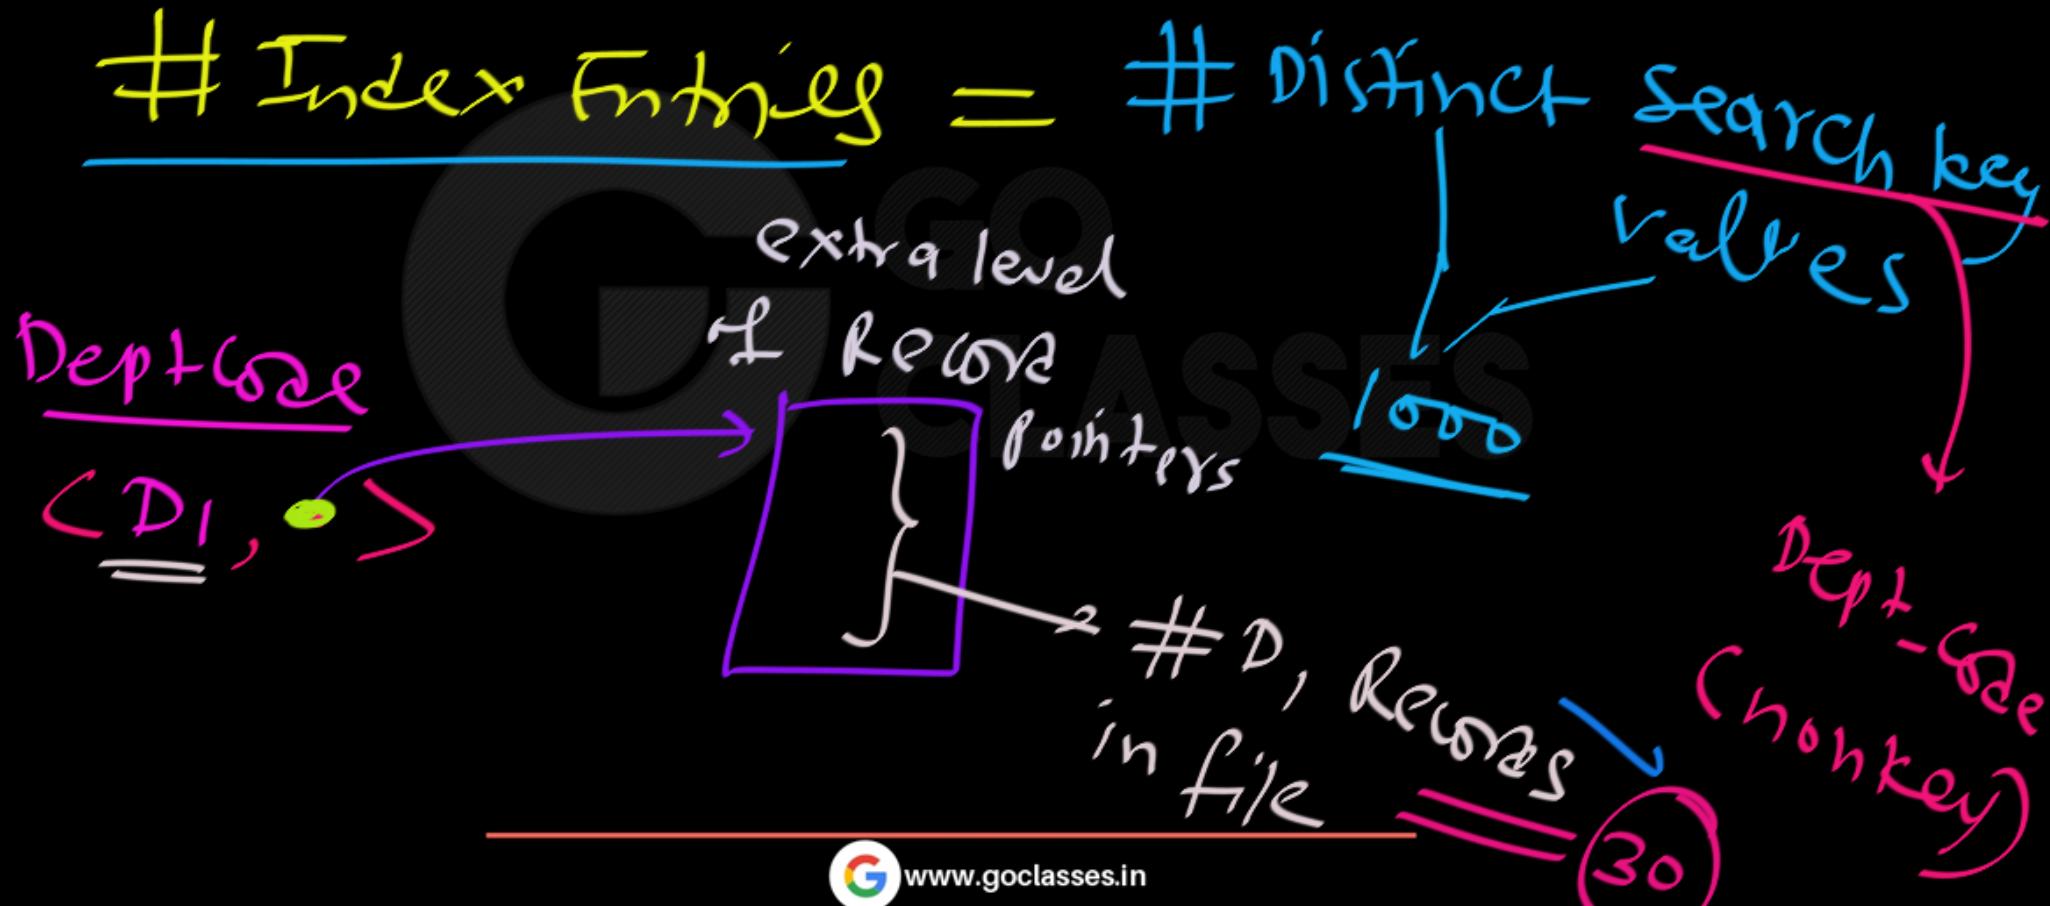
PI on ILI



SI on Non-key:



1st level Index :

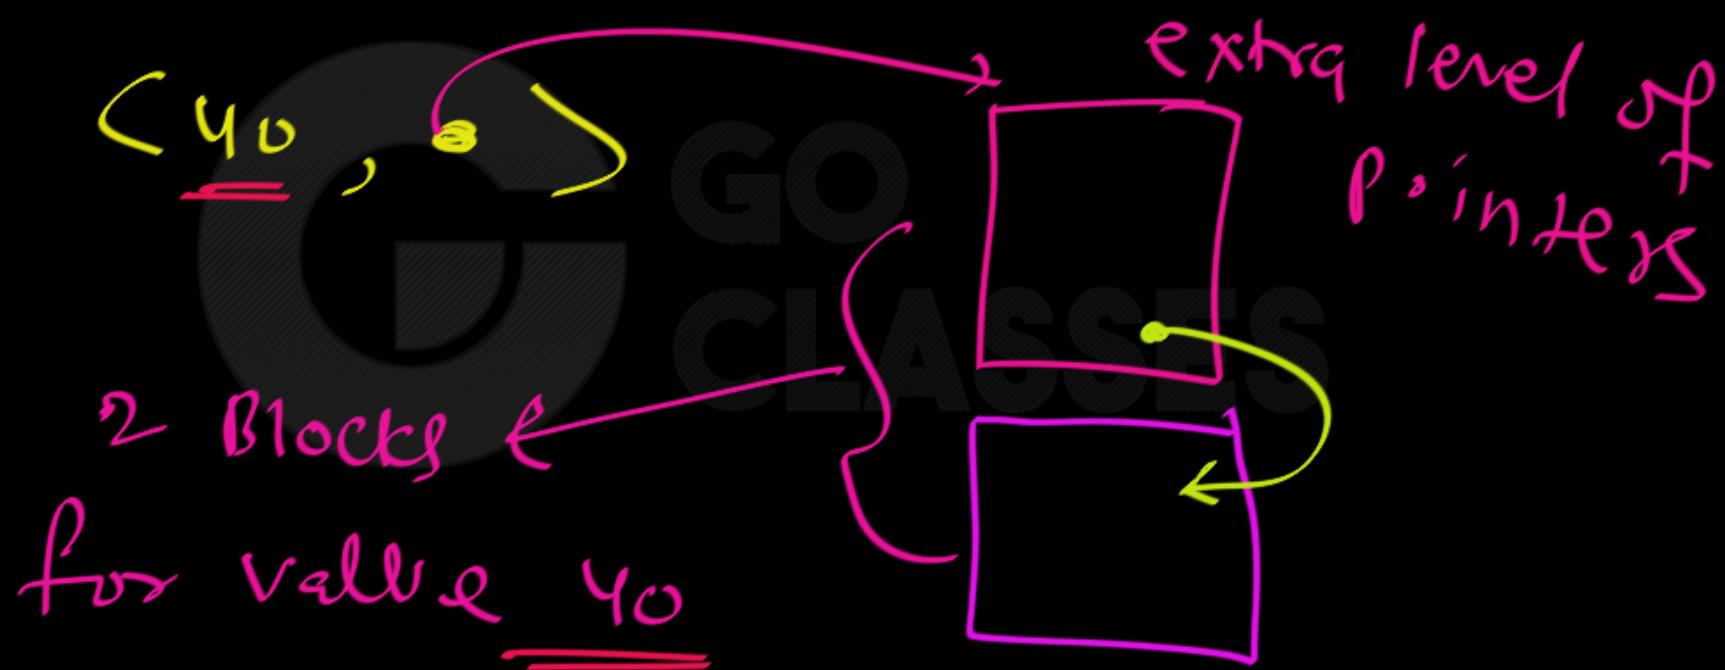


Can we put 30 Record points in a single Block ? $\Rightarrow \underline{\text{Yes}}$ ✓

$$\text{max } \#RP \text{ Per Block} = \left\lfloor \frac{512 \text{ B}}{7 \text{ B}} \right\rfloor = 73$$

By chance :

SI on Non key :



Coming Back to Question:

#Blocks in the Extra level of
indirection = 1000 ✓

1st level of Index: IL1:

#Index entries = 1000

IL1 entry size = $9B + \underline{6B} = 15B$

for every level
of Index:

Index entry size = 15B

Blocking factor = $\left\lfloor \frac{512}{15} \right\rfloor = 34$

Block
pointer
size

for Dept - Code size

$$\frac{1^{\text{st}} \text{ level index blocks}}{1^{\text{st}} \text{ level index blocks}} = \left\lceil \frac{1000}{34} \right\rceil = 30$$

$$\frac{2^{\text{nd}} \text{ level } \# \text{ entries}}{2^{\text{nd}} \text{ level } \# \text{ entries}} = 30$$

$$\frac{2^{\text{nd}} \text{ level } \# \text{ blocks}}{2^{\text{nd}} \text{ level } \# \text{ blocks}} = \left\lceil \frac{30}{34} \right\rceil = 1$$

levels = $2 + 1 \Rightarrow$ levels of index
extra level for indirection

(i) 34 (ii) 1000 (iii) 1000 ; 30

(iv) $2 + 1$ unorderes

(v) extra level of indirection

$$1 + 30 + 1000 = 1031$$

$$\frac{2+1}{33} + \text{(worst case)}$$

extra level of indirection

Q 17.18: FUNDAMENTALS OF Database Systems 7th Ed, Navathe, Elmasri

- 17.18. Consider a disk with block size $B = 512$ bytes. A block pointer is $P = 6$ bytes long, and a record pointer is $P_R = 7$ bytes long. A file has $r = 30,000$ EMPLOYEE records of *fixed length*. Each record has the following fields: Name (30 bytes), Ssn (9 bytes), Department_code (9 bytes), Address (40 bytes), Phone (10 bytes), Birth_date (8 bytes), Sex (1 byte), Job_code (4 bytes), and Salary (4 bytes, real number). An additional byte is used as a deletion marker.

CLASSES
Part f: On Next Page

- f. Suppose that the file is *ordered* by the nonkey field Department_code and we want to construct a clustering index on Department_code that uses block anchors (every new value of Department_code starts at the beginning of a new block). Assume there are 1,000 distinct values of Department_code and that the EMPLOYEE records are evenly distributed among these values. Calculate (i) the index blocking factor bfr_i (which is also the index fan-out fo); (ii) the number of first-level index entries and the number of first-level index blocks; (iii) the number of levels needed if we make it into a multi-level index; (iv) the total number of blocks required by the multilevel index; and (v) the number of block accesses needed to search for and retrieve all records in the file that have a specific Department_code value, using the clustering index (assume that multiple blocks in a cluster are contiguous).

Multi-level Secondary Indexes

Secondary indexes can also be converted to multi-level indexes

First level index

- as many entries as there are records in the data file



First level index is an ordered file

so, in the second level index, the number of entries will be equal to the number of *blocks* in the first level index rather than the number of *records*

Similarly in other higher levels

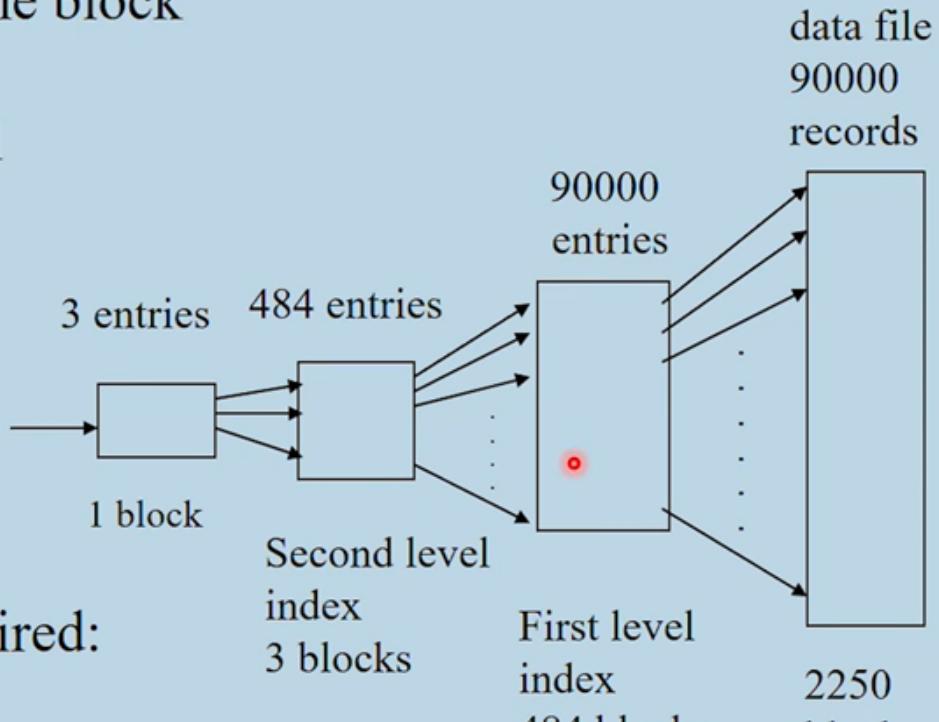
Making the Secondary Index Multi-level

Multilevel Index –

Successive levels of indices are built
till the last level has one block

height – no. of levels

block accesses: height + 1



For the example data file:

No of block accesses required:

multi-level index: 4

single level index: 10

A multilevel index considers the index file, which we will now refer to as the **first (or base) level** of a multilevel index, as an *ordered file* with a *distinct value* for each $K(i)$.

Therefore, by considering the first-level index file as a sorted data file, we can create a primary index for the first level; this index to the first level is called the **second level** of the multilevel index. Because the second level is a primary index, we can use block anchors so that the second level has one entry for *each block* of the first level. The blocking factor bfr_i for the second level—and for all subsequent levels—is the same as that for the first-level index because all index entries are the same size; each has one field value and one block address. If the first level has r_1 entries, and the blocking factor—which is also the fan-out—for the index is $bfr_i = fo$, then the first level needs $\lceil(r_1/fo)\rceil$ blocks, which is therefore the number of entries r_2 needed at the second level of the index.

We can repeat this process for the second level. The **third level**, which is a primary index for the second level, has an entry for each second-level block, so the number of third-level entries is $r_3 = \lceil (r_2/fo) \rceil$. Notice that we require a second level only if the first level needs more than one block of disk storage, and, similarly, we require a third level only if the second level needs more than one block. We can repeat the preceding process until all the entries of some index level t fit in a single block. This block at the t th level is called the **top** index level.⁶ Each level reduces the number of entries at the previous level by a factor of fo —the index fan-out—so we can use the formula $1 \leq (r_1/((fo)^t))$ to calculate t . Hence, a multilevel index with r_1 first-level entries will have approximately t levels, where $t = \lceil (\log_{fo}(r_1)) \rceil$. When searching the index, a single disk block is retrieved at each level. Hence, t disk blocks are accessed for an index search, where t is the *number of index levels*.

The multilevel scheme described here can be used on any type of index—whether it is primary, clustering, or secondary—as long as the first-level index has *distinct values for $K(i)$ and fixed-length entries*. Figure 17.6 shows a multilevel index built over



Next Topic:

ISAM

Indexed sequential
access method

I SAM:

multilevel

PI

by Default

ISAM

Sequential file
ordered by key

ISAM

I SAM:

multilevel
CT + sequential file
ordered by Non key

I SAM

A common file organization used in business data processing is an ordered file with a multilevel primary index on its ordering key field. Such an organization is called an indexed sequential file and was used in a large number of early IBM systems.

I SAM

I sAM file : orders file with

multilevel primary
Index

A common file organization used in business data processing is an ordered file with a multilevel primary index on its ordering key field. Such an organization is called an indexed sequential file and was used in a large number of early IBM systems.

I SAM

I sAm file :

A file with
multilevel primary
Index

Index Sequential Access Method (ISAM) Files

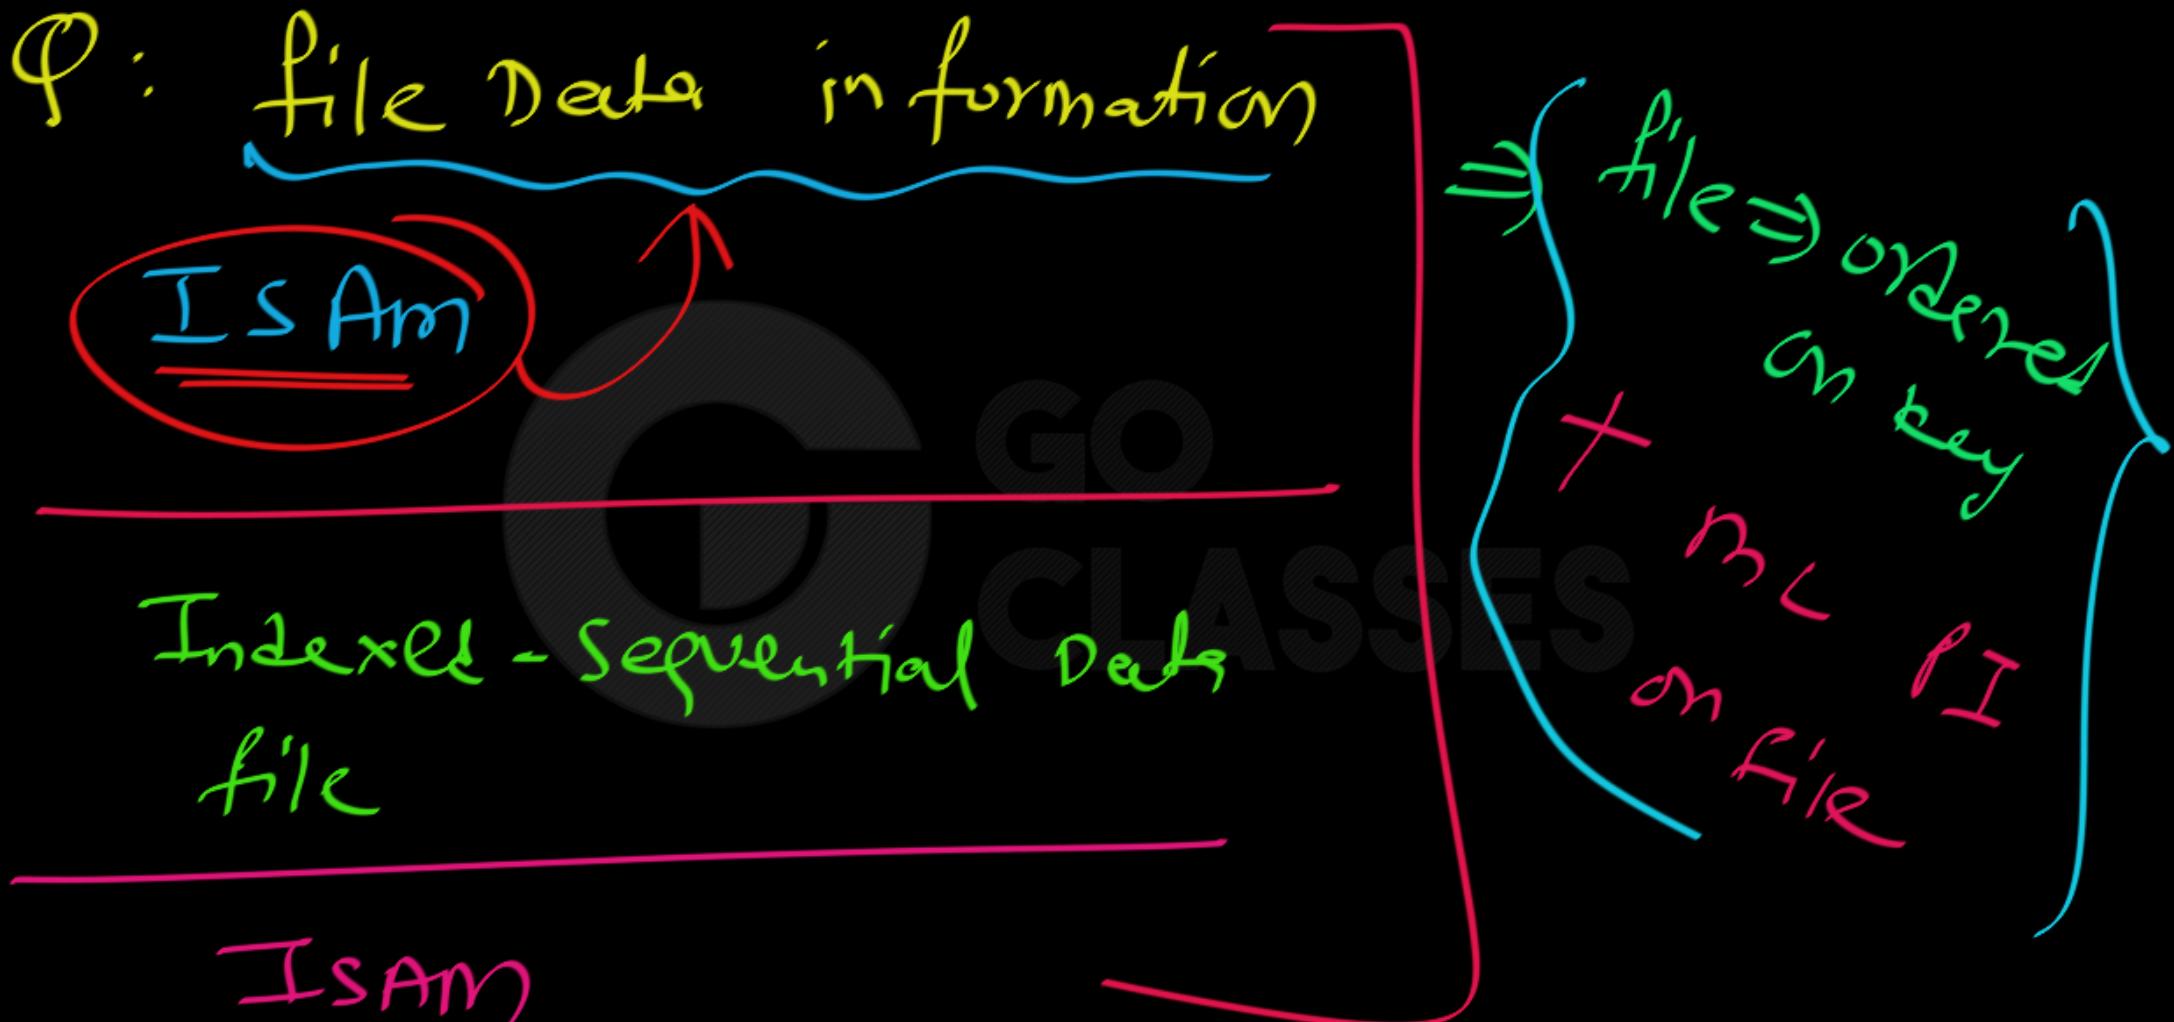
ISAM files –

Ordered files with a multilevel primary/clustering index

Most suitable for files that are relatively static

If the files are dynamic, we need to go for dynamic multi-level index structures based on B^+ - trees

NPTEL IITM



I SAM

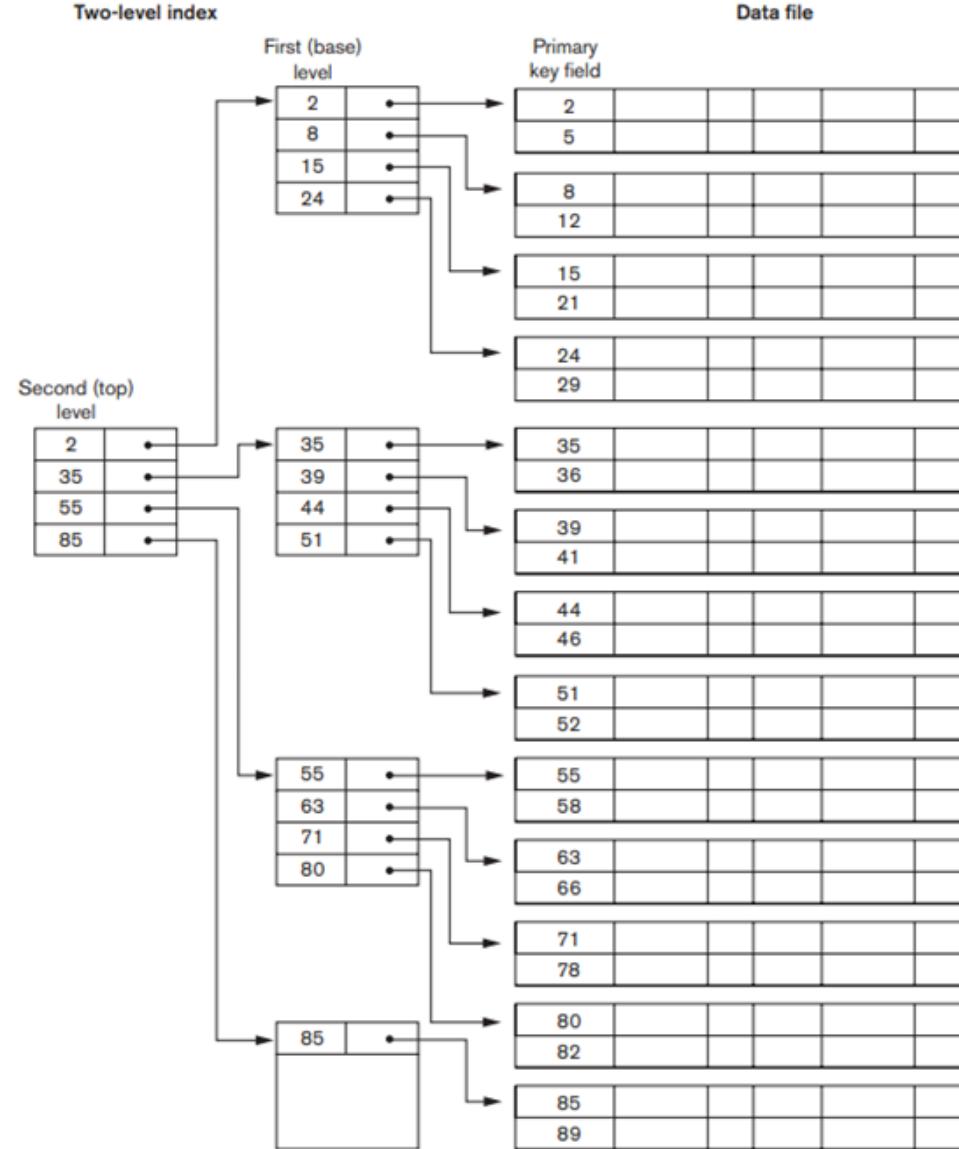
= multilevel primary/clustered

Index

= Index sequential file

Figure 17.6

A two-level primary index resembling ISAM (indexed sequential access method) organization.



Multilevel SI ; 1st level : SI

~~Remaining level : Sparse~~

k^{th} level is PI

on $k-1$ level

Multilevel CI :- 1st level : CI

~~Remaining level : Sparse~~

kth level is

PI

on k-1 level

Multilevel PI

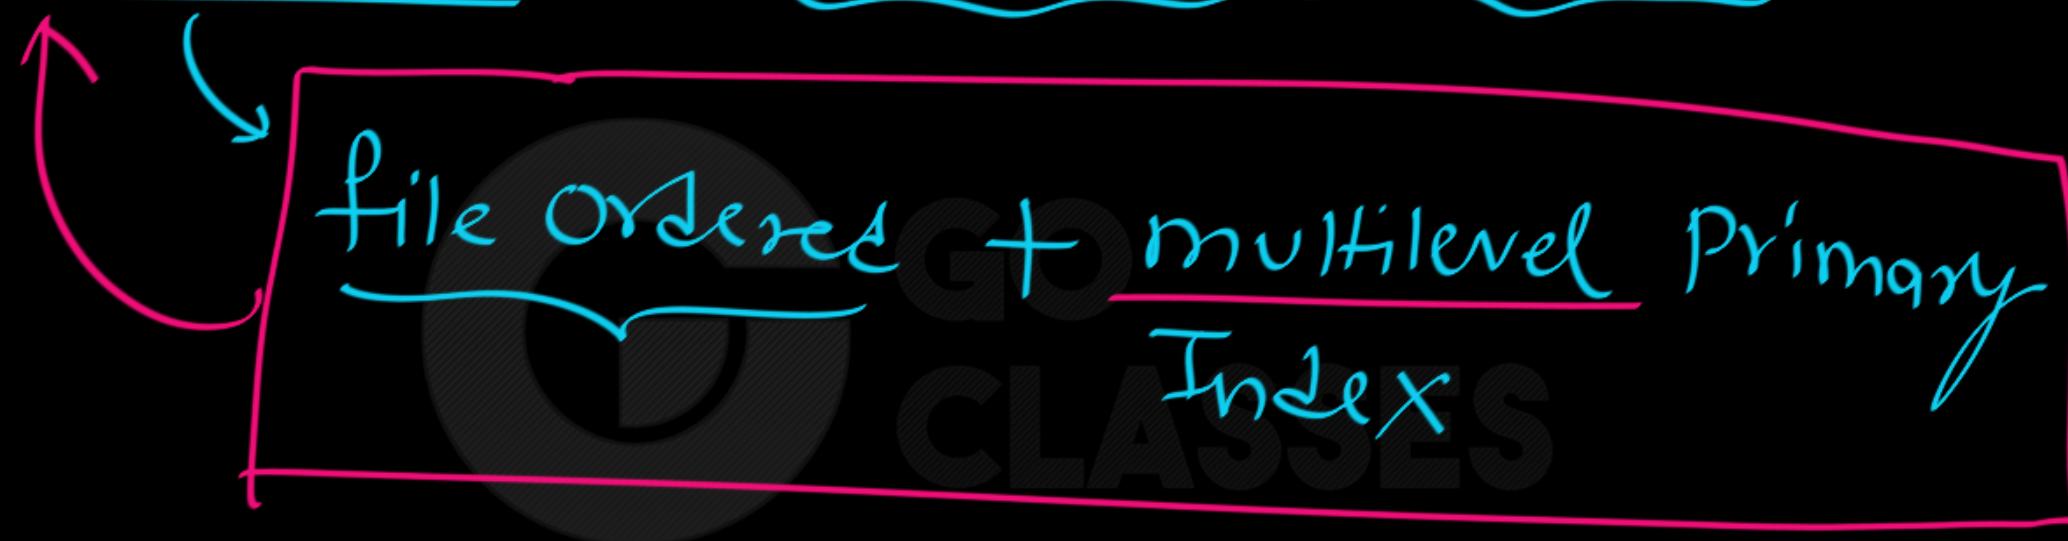
1st level : PI

Remaining level : Sparse

kth level is PI

on k-1 level

ISAM file = Indexed sequential file



3.7.3 Indexing: GATE CSE 1993 | Question: 14 top ↴<https://gateoverflow.in/2311>

An ISAM (indexed sequential) file consists of records of size 64 bytes each, including key field of size 14 bytes. An address of a disk block takes 2 bytes. If the disk block size is 512 bytes and there are $16K$ records, compute the size of the data and index areas in terms of number blocks. How many levels of tree do you have for the index?

multilevel primary index

3.7.3 Indexing: GATE CSE 1993 | Question: 14 top ↗

↗ <https://gateoverflow.in/2311>



An ISAM (indexed sequential) file consists of records of size 64 bytes each, including key field of size 14 bytes. An address of a disk block takes 2 bytes. If the disk block size is 512 bytes and there are 16K records, compute the size of the data and index areas in terms of number blocks. How many levels of tree do you have for the index?

multilevel
primary index

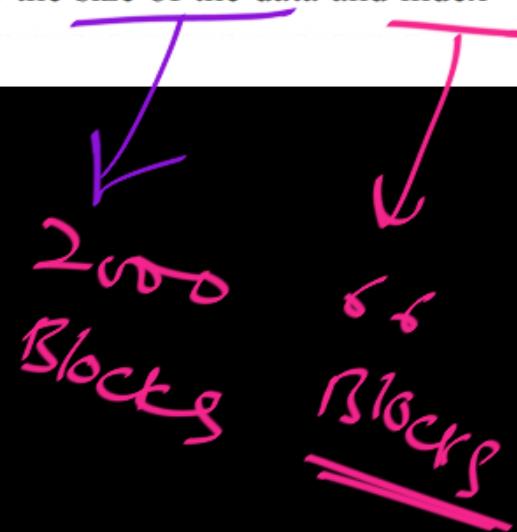
ISAM file

ison

3.7.3 Indexing: GATE CSE 1993 | Question: 14 top ↗<https://gateoverflow.in/2311>

An ISAM (indexed sequential) file consists of records of size 64 bytes each, including key field of size 14 bytes. An address of a disk block takes 2 bytes. If the disk block size is 512 bytes and there are 16K records, compute the size of the data and index areas in terms of number blocks. How many levels of tree do you have for the index?

3 levels



file : Record size = 64 B

key field = 14 B

Block address (BP) = 2 B

Block size = 512 B

Records = 16000

file Blocks

$$\lceil \frac{16000}{8} \rceil$$

= 2000
Blocks

→ Blocking factor of file = #Records / Block

$$\left\lfloor \frac{512}{64} \right\rfloor = 8$$

Index:

multi level primary index

Index entry size:

(whatever level
of index)

key field
size

$$14 + 2B = \underline{16B}$$

I SAM file

Index Blocking factor :

$$\left\lfloor \frac{512B}{16B} \right\rfloor = 32 \text{ Index entries per block}$$

1st level index

$$\# \text{index entries} = \# \text{Blocks in file}$$

$$\# IL_1 \text{ Blocks} = \left\lceil \frac{2000}{32} \right\rceil = 63 \text{ } IL_1 \text{ } \underline{\text{Blocks}}$$

2ⁿ level index:

#IL2 entries = #Blocks in IL1

= 63

$$\#IL2 \text{ Blocks} = \left\lceil \frac{63}{32} \right\rceil = 2$$

3ⁿ level index: → 1 Block; 2 entries

Size of Data file: 2000 Blocks

Size of Index file: 1 + 2 + 63 = 66 Blocks

ISAM file

multilevel primary index
master index

3.7.2 Indexing: GATE CSE 1990 | Question: 10b top

<https://gateoverflow.in/85691>



One giga bytes of data are to be organized as an indexed-sequential file with a uniform blocking factor 8. Assuming a block size of 1 Kilo bytes and a block referencing pointer size of 32 bits, find out the number of levels of indexing that would be required and the size of the index at each level. Determine also the size of the master index. The referencing capability (fanout ratio) per block of index storage may be considered to be 32.

file blocking factor

file with
multi-level primary
index

3.7.2 Indexing: GATE CSE 1990 | Question: 10b top

<https://gateoverflow.in/85691>



One giga bytes of data are to be organized as an indexed-sequential file with a uniform blocking factor 8. Assuming a block size of 1 Kilo bytes and a block referencing pointer size of 32 bits, find out the number of levels of indexing that would be required and the size of the index at each level. Determine also the size of the master index. The referencing capability (fanout ratio) per block of index storage may be considered to be 32.

#index entries = 32

Total Index size

subjective

Total index size = $2^{15} + 2^{10} + 2^5 + 1$ Blocks

3.7.2 Indexing: GATE CSE 1990 | Question: 10b top

<https://gateoverflow.in/85691>



One giga bytes of data are to be organized as an indexed-sequential file with a uniform blocking factor 8. Assuming a block size of 1 Kilo bytes and a block referencing pointer size of 32 bits, find out the number of levels of indexing that would be required and the size of the index at each level. Determine also the size of the master index. The referencing capability (fanout ratio) per block of index storage may be considered to be 32.

1st level

2^{15} blocks

2nd

2^{10}

3rd

2^5

4th

1 block

Data file: 1 GB \Rightarrow

Blocking factor: 8 Records / Block

Block size = 1 kB

Block pointer = 4B

$$\# \text{Blocks} = \frac{1 \text{ GB}}{1 \text{ kB}} = 1m = 2^{20} \text{ blocks}$$

$$\text{Record size} = \frac{1 \text{ kB}}{8} = 128 \text{ B}$$

1st level index:

$$\# \text{Entries} = 2^{20} = \# \text{file blocks}$$

Blocking factor of index = 32

$$\# \text{IL1 Blocks} = \frac{2^{20}}{2^5} = 2^{15} \text{ Blocks}$$

2ⁿ⁻¹ level index:

$$\# \text{Entries} = 2^{15} = \# \text{IL1 Blocks}$$

Blocking factor of index = 32

$$\# \text{IL1 Blocks} = \left\lceil \frac{2^{15}}{2^5} \right\rceil = 2 \text{ Blocks}$$

3rd level index:

$$\# \text{Entries} = 2^{10} = \# \text{IL2 Blocks}$$

Blocking factor of index = 32

$$\# \text{IL1 Blocks} = \left\lceil \frac{2^{10}}{2^5} \right\rceil = \underline{\underline{2^5 \text{ Blocks}}}$$

4th level index:

$$\# \text{Entries} = 2^5 = \# \text{IL2 Blocks}$$

Blocking factor of index = 32

$$\# \text{IL1 Blocks} = \left\lceil \frac{2^5}{2^5} \right\rceil = 1 \text{ Blocks}$$

Index levels = 4 ✓

3.7.6 Indexing: GATE CSE 2008 | Question: 70 top ↴<https://gateoverflow.in/259>

Consider a file of 16384 records. Each record is 32 *bytes* long and its key field is of size 6 *bytes*. The file is ordered on a non-key field, and the file organization is unspanned. The file is stored in a file system with block size 1024 *bytes*, and the size of a block pointer is 10 *bytes*. If the secondary index is built on the key field of the file, and a multi-level index scheme is used to store the secondary index, the number of first-level and second-level blocks in the multi-level index are respectively

- A. 8 and 0
- B. 128 and 6
- C. 256 and 4
- D. 512 and 5

[tests.gatecse.in](#)

gate2008-cse

databases

indexing

normal

[goclasses.in](#)[tests.gatecse.in](#)

www.goclasses.in

3.7.6 Indexing: GATE CSE 2008 | Question: 70 top ↴<https://gateoverflow.in/259>

Consider a file of 16384 records. Each record is 32 bytes long and its key field is of size 6 bytes. The file is ordered on a non-key field, and the file organization is unspanned. The file is stored in a file system with block size 1024 bytes, and the size of a block pointer is 10 bytes. If the secondary index is built on the key field of the file, and a multi-level index scheme is used to store the secondary index, the number of first-level and second-level blocks in the multi-level index are respectively

- A. 8 and 0
- B. 128 and 6
- C. 256 and 4
- D. 512 and 5

gate2008-cse databases indexing normal

tests.gatecse.in

goclasses.in

file is ordered by non-key field ✓
Index on key field : SI on key

search key of Index

file:

16384 records.

Record size = 32 B

key field = 6B

Block size = 1024 B

" pointer = 10 B

SI on key ;

Index entry size : $6B + 10B = 16B$

Index Blocking factor = $\left\lceil \frac{1024B}{16B} \right\rceil = 64$

1st level index:

$$\# \text{ ILI entries} = 16384$$

$$\# \downarrow \text{Blocks}_{\text{ILI}} = \left\lceil \frac{16384}{64} \right\rceil = 256$$

256 blocks on 1st level,

2ⁿ⁻¹ level index.

$$\# \text{ IL2 entries} = 256$$

$$\# \downarrow \text{Blocks} = \left\lceil \frac{256}{64} \right\rceil = 4$$

4 blocks in level 2

3rd level index.

$$\# \text{ IL3 entries} = 9$$

$$\# \downarrow \text{Blocks} = \left\lceil \frac{9}{64} \right\rceil = 1$$

1 blocks
in level 3

Next Topic:

Indexing

Recap (Self-Read)

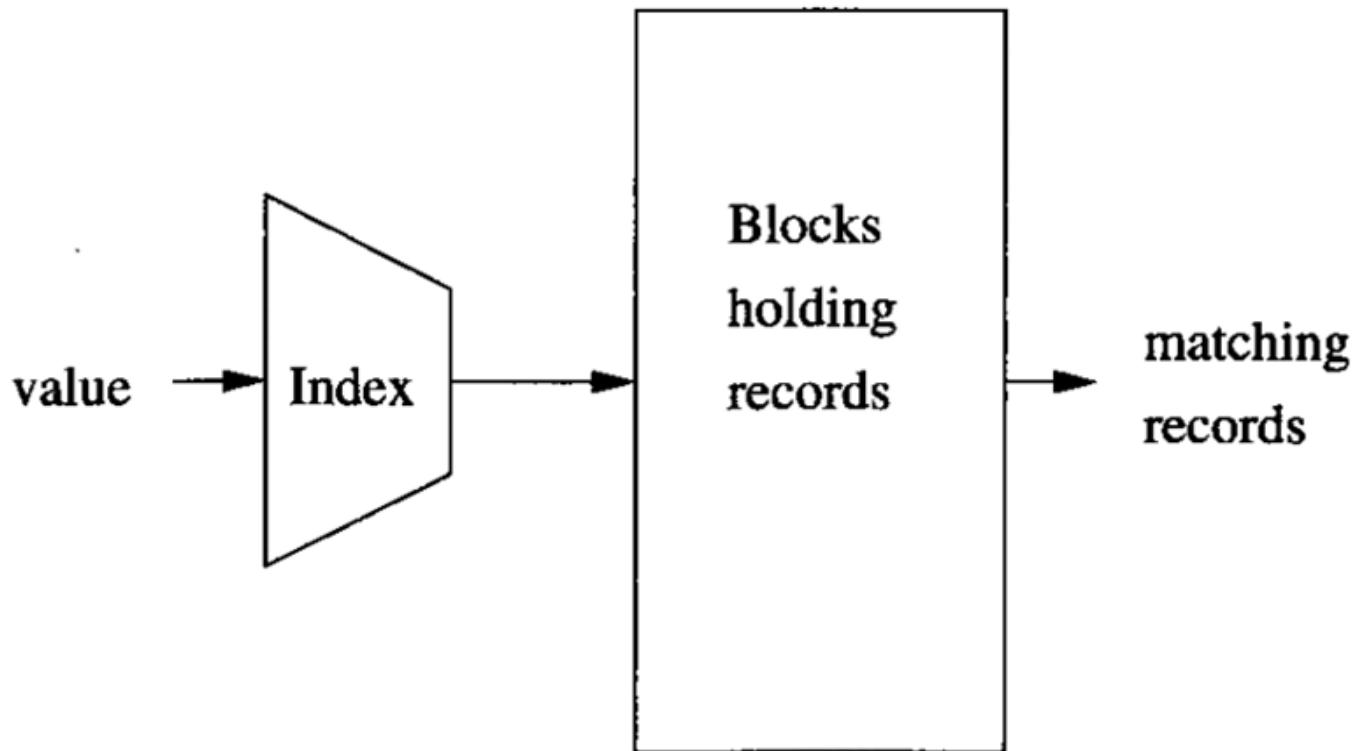


Figure 14.1: An index takes a value for some field(s) and finds records with the matching value

Many queries reference only a small proportion of the records in a file. For example, a query like “Find all instructors in the Physics department” or “Find the total number of credits earned by the student with *ID* 22201” references only a fraction of the student records. It is inefficient for the system to read every tuple in the *instructor* relation to check if the *dept_name* value is “Physics”. Likewise, it is inefficient to read the entire *student* relation just to find the one tuple for the *ID* “32556.”. Ideally, the system should be able to locate these records directly. To allow these forms of access, we design additional structures that we associate with files.

In this chapter, we assume that a file already exists with some primary organization such as the unordered, ordered organizations that were described in last chapter.

We will describe **additional auxiliary access structures called indexes**, which are used to speed up the retrieval of records in response to certain search conditions. The index structures are additional files on disk.

Indexes enable efficient access to records based on the indexing fields that are used to construct the index.

Basically, any field of the file can be used to create an index, and multiple indexes on different fields can be constructed on the same file.

To find a record or records in the data file based on a search condition on an indexing field, the index is searched, which leads to pointers to one or more disk blocks in the data file where the required records are located.

Different Kinds of “Keys”

There are many meanings of the term “key.” We used it in Section 2.3.6 to mean the primary key of a relation. We shall also speak of “sort keys,” the attribute(s) on which a file of records is sorted. We just introduced “search keys,” the attribute(s) for which we are given values and asked to search, through an index, for tuples with matching values. We try to use the appropriate adjective — “primary,” “sort,” or “search” — when the meaning of “key” is unclear.

We often want to have more than one index for a file. For example, we may wish to search for a book by author, by subject, or by title.

An attribute or set of attributes used to look up records in a file is called a **search key**. Note that this definition of *key* differs from that used in *primary key*, *candidate key*, and *superkey*. This duplicate meaning for *key* is (unfortunately) well

For a file with a given record structure consisting of several fields (or attributes), an index access structure is usually defined on a single field of a file, called an **indexing field** (or **indexing attribute**).¹ The index typically stores each value of the index field along with a list of pointers to all disk blocks that contain records with that field value. The values in the index are *ordered* so that we can do a *binary search* on the index. If both the data file and the index file are ordered, and since the index file is typically much smaller than the data file, searching the index using a binary search is a better option. Tree-structured multilevel indexes (see Section 17.2) implement

We describe different types of single-level ordered indexes—primary, secondary, and clustering—in Section 17.1. By viewing a single-level index as an ordered file, one can develop additional indexes for it, giving rise to the concept of multilevel indexes. A popular indexing scheme called **ISAM (indexed sequential access method)** is based on this idea. We discuss multilevel tree-structured indexes in Sec-

There are several types of ordered indexes. A **primary index** is specified on the *ordering key field* of an **ordered file** of records. Recall from Section 16.7 that an ordering key field is used to *physically order* the file records on disk, and every record has a *unique value* for that field. If the ordering field is not a key field—that is, if numerous records in the file can have the same value for the ordering field—another type of index, called a **clustering index**, can be used. The data file is called a **clustered file** in this latter case. Notice that a file can have at most one physical ordering field, so it can have at most one primary index or one clustering index, *but not both*.

A third type of index, called a **secondary index**, can be specified on any *nonordering* field of a file. A data file can have several secondary indexes in addition to its primary access method. We discuss these types of single-level indexes in the next three subsections.

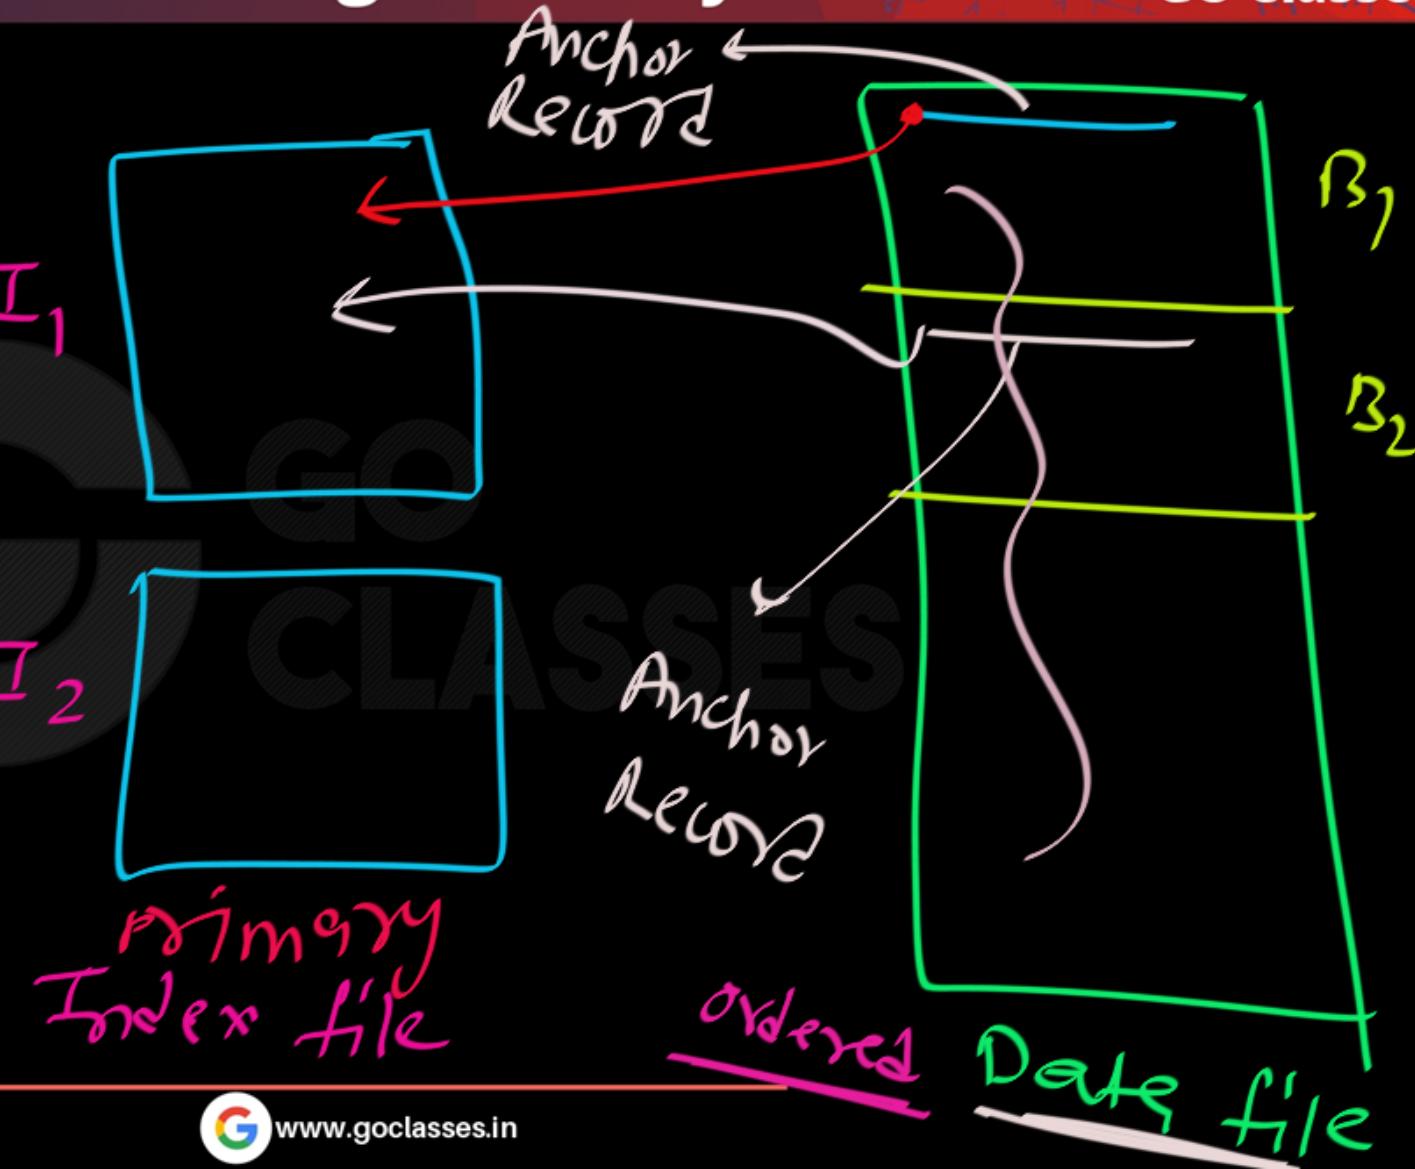
Ordered Indexes:

Problem (DisAdvantage)

modification is very time consuming (Very Hard)

+ insertion
+ deletion
+ modify

Inserting
a
Record
that should
come at
the end
of B_1
In all blocks
new
Anchor record

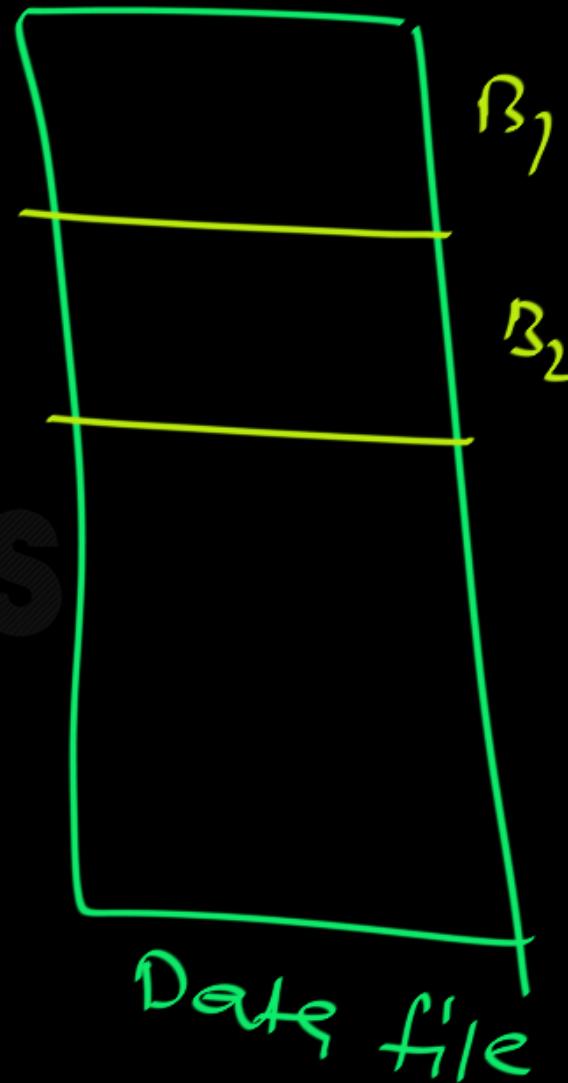


on insertion

almost entire
index needs
to be
changed



Dense
Index file



Dense file

If Multilevel index ; modification becomes even more hairy.

When to prefer these indexes that we have seen

so far ? \Rightarrow

Static indexes

When modifications are rare .
Data is mostly static .

for Dynamically changing Data file,

we need some "modification-friendly"

Index

Dynamic
Indexes

B⁺ Tree Index

B Tree index

better

\mathcal{B} , $\mathcal{B}+$ Tree Indexing

GO
CLASSES
Next...